

Systems and Networking – Unit I

B.Sc. in Applied Computer Science and Artificial Intelligence

2022-2023

Gabriele Tolomei

Department of Computer Science

Sapienza Università di Roma

tolomei@di.uniroma1.it



SAPIENZA
UNIVERSITÀ DI ROMA

Recap from Last Lecture: Synchronization

- Concurrent accesses to shared resources by multiple cooperating processes/threads can lead to unexpected or erroneous behavior
- Process/Thread cooperation must guarantee consistency of any shared data/resource, **regardless of CPU scheduling**
- Maintaining shared data consistency requires mechanisms to ensure synchronized execution of critical sections by processes/threads
- Critical sections are specific pieces of code which contain shared resources that need to be "protected"

Recap from Last Lecture: Locks

- Synchronization primitives ensure that only one process/thread at a time executes in a critical section (**mutual exclusion**)
- Locks allow protection of critical sections by atomically testing and taking/releasing the access to a critical section
- Locks can be implemented leveraging some HW support:
 - **disabling interrupts** (can miss or delay important events)
 - **atomic instructions** (busy waiting/spinlock inefficient)

Higher-Level Synchronization Primitives

- More general synchronization mechanisms
 - Not only for safely accessing critical sections

Higher-Level Synchronization Primitives

- More general synchronization mechanisms
 - Not only for safely accessing critical sections
- 2 common high-level synchronization primitives:
 - **Semaphores:** binary (mutex) and counting
 - **Monitors:** mutex and condition variables

Semaphores: Overview

- Another data structure that provides mutual exclusion to critical sections

Semaphores: Overview

- Another data structure that provides mutual exclusion to critical sections
- Can also play the role of an atomic counter

Semaphores: Overview

- Another data structure that provides mutual exclusion to critical sections
- Can also play the role of an atomic counter
- Generalization of locks invented by **Dijkstra** in 1965

Semaphores: Overview

- Another data structure that provides mutual exclusion to critical sections
- Can also play the role of an atomic counter
- Generalization of locks invented by **Dijkstra** in 1965
- Special type of (integer) variable that supports **2 atomic operations**
 - **wait()** (also **P()**): decrement, block until semaphore is open
 - **signal()** (also **V()**): increment, allow another thread to enter

Blocking in Semaphores

- Associated with each semaphore is a queue of waiting processes/threads

Blocking in Semaphores

- Associated with each semaphore is a queue of waiting processes/threads
- When **wait()** is called by a thread:
 - If semaphore is open thread continues, otherwise thread blocks on queue

Blocking in Semaphores

- Associated with each semaphore is a queue of waiting processes/threads
- When **wait()** is called by a thread:
 - If semaphore is open thread continues, otherwise thread blocks on queue
- Then **signal()** opens the semaphore:
 - If a thread is waiting on the queue the thread is unblocked, whilst if no threads are waiting on the queue, the signal is remembered for the next thread

Blocking in Semaphores

- Associated with each semaphore is a queue of waiting processes/threads
- When **wait()** is called by a thread:
 - If semaphore is open thread continues, otherwise thread blocks on queue
- Then **signal()** opens the semaphore:
 - If a thread is waiting on the queue the thread is unblocked, whilst if no threads are waiting on the queue, the signal is remembered for the next thread
- In other words, **signal()** is stateful and has "history"

Semaphores: Types

- **Binary Semaphore** a.k.a. Mutex (same as a Lock)
 - Guarantees mutually exclusive access to a resource (i.e., only one process/thread executes in a critical section)
 - Its associated integer variable can only take 2 values: 0 or 1
 - Initialized to open (e.g., `value = 1`)

Semaphores: Types

- **Binary Semaphore** a.k.a. Mutex (same as a Lock)
 - Guarantees mutually exclusive access to a resource (i.e., only one process/thread executes in a critical section)
 - Its associated integer variable can only take 2 values: 0 or 1
 - Initialized to open (e.g., `value = 1`)
- **Counting Semaphore**
 - To manage multiple shared resources
 - The initial value of the semaphore is usually the number of resources
 - A process can access to a resource as long as at least one is available

Semaphores: Key Ideas

```
// Semaphore S  
  
S.wait(); // wait until S is available  
  
<critical section>  
  
S.signal(); notify other processes that S is open
```


Semaphores: Key Ideas

```
// Semaphore S  
  
S.wait(); // wait until S is available  
  
<critical section>  
  
S.signal(); notify other processes that S is open
```

Each semaphore supports a queue of processes that are waiting to access the critical section (e.g., to buy milk)

Semaphores: Key Ideas

```
// Semaphore S  
  
S.wait(); // wait until S is available  
  
<critical section>  
  
S.signal(); notify other processes that S is open
```

Each semaphore supports a queue of processes that are waiting to access the critical section (e.g., to buy milk)

If a process executes **S.wait()** and semaphore **S** is open (non-zero), it continues executing, otherwise the OS puts the process on the wait queue

Semaphores: Key Ideas

```
// Semaphore S  
  
S.wait(); // wait until S is available  
  
<critical section>  
  
S.signal(); notify other processes that S is open
```

Each semaphore supports a queue of processes that are waiting to access the critical section (e.g., to buy milk)

If a process executes **S.wait()** and semaphore **S** is open (non-zero), it continues executing, otherwise the OS puts the process on the wait queue

A **S.signal()** unblocks one process on semaphore **S**'s wait queue

Binary Semaphore: Example

"Too Much Milk" Using **Lock**

```
# Thread Bob  
  
Lock.acquire()  
  
if (!milk):  
    buy_milk()  
  
Lock.release()
```

```
# Thread Carla  
  
Lock.acquire()  
  
if (!milk):  
    buy_milk()  
  
Lock.release()
```

Binary Semaphore: Example

"Too Much Milk" Using **Lock**

```
# Thread Bob  
  
Lock.acquire()  
  
if (!milk):  
    buy_milk()  
  
Lock.release()
```

```
# Thread Carla  
  
Lock.acquire()  
  
if (!milk):  
    buy_milk()  
  
Lock.release()
```

"Too Much Milk" Using **Semaphore**

```
# Thread Bob  
  
S.wait()  
  
if (!milk):  
    buy_milk()  
  
S.signal()
```

```
# Thread Carla  
  
S.wait()  
  
if (!milk):  
    buy_milk()  
  
S.signal()
```

Binary Semaphore: Example

"Too Much Milk" Using **Lock**

"Too Much Milk" Using **Semaphore**

# Thread Bob	# Thread Carla	# Thread Bob	# Thread Carla
<code>Lock.acquire()</code>	<code>Lock.acquire()</code>	<code>S.wait()</code>	<code>S.wait()</code>
<code>if (!milk):</code> <code> buy_milk()</code>	<code>if (!milk):</code> <code> buy_milk()</code>	<code>if (!milk):</code> <code> buy_milk()</code>	<code>if (!milk):</code> <code> buy_milk()</code>
<code>Lock.release()</code>	<code>Lock.release()</code>	<code>S.signal()</code>	<code>S.signal()</code>

Semaphore: Implementation

```
Class Semaphore {  
    public void wait(Thread t);  
    public void signal();  
    private int value;  
    private int guard;  
    private Queue q;  
  
    Semaphore(int val) {  
        // initialize semaphore  
        // with val and empty queue  
        this.value = val;  
        this.q = null;  
    }  
}
```

Semaphore: Implementation

```
Class Semaphore {  
    public void wait(Thread t);  
    public void signal();  
    private int value;  
    private int guard;  
    private Queue q;  
  
    Semaphore(int val) {  
        // initialize semaphore  
        // with val and empty queue  
        this.value = val;  
        this.q = null;  
    }  
}
```

```
public void wait(Thread t) {  
    while(test&set(this.guard) == 1) {  
        // while busy do nothing  
    }  
    this.value -= 1;  
    if(this.value < 0) {  
        q.push(t);  
        t.sleep_and_reset_guard_to_0();  
    }  
    else {  
        this.guard = 0;  
    }  
}
```


Semaphore: Implementation

```
Class Semaphore {  
    public void wait(Thread t);  
    public void signal();  
    private int value;  
    private int guard;  
    private Queue q;  
  
    Semaphore(int val) {  
        // initialize semaphore  
        // with val and empty queue  
        this.value = val;  
        this.q = null;  
    }  
}
```

```
public void wait(Thread t) {  
    while(test&set(this.guard) == 1) {  
        // while busy do nothing  
    }  
    this.value -= 1;  
    if(this.value < 0) {  
        q.push(t);  
        t.sleep_and_reset_guard_to_0();  
    }  
    else {  
        this.guard = 0;  
    }  
}
```

```
public void signal() {  
    while(test&set(this.guard) == 1) {  
        // while busy do nothing  
    }  
    this.value += 1;  
    if(!q.isEmpty()) {  
        t = q.pop();  
        push_onto_ready_queue(t);  
    }  
    this.guard = 0;  
}
```

Semaphore: Implementation

```
Class Semaphore {  
    public void wait(Thread t);  
    public void signal();  
    private int value;  
    private int guard;  
    private Queue q;  
  
    Semaphore(int val) {  
        // initialize semaphore  
        // with val and empty queue  
        this.value = val;  
        this.q = null;  
    }  
}
```

```
public void wait(Thread t) {  
    while(test&set(this.guard) == 1) {  
        // while busy do nothing  
    }  
    this.value -= 1;  
    if(this.value < 0) {  
        q.push(t);  
        t.sleep_and_reset_guard_to_0();  
    }  
    else {  
        this.guard = 0;  
    }  
}
```

```
public void signal() {  
    while(test&set(this.guard) == 1) {  
        // while busy do nothing  
    }  
    this.value += 1;  
    if(!q.isEmpty()) {// this.value <= 0  
        t = q.pop();  
        push_onto_ready_queue(t);  
    }  
    this.guard = 0;  
}
```

Semaphore: Implementation

```
Class Semaphore {  
    public void wait(Thread t);  
    public void signal();  
    private int value;  
    private int guard;  
    private Queue q;  
  
    Semaphore(int val) {  
        // initialize semaphore  
        // with val and empty queue  
        this.value = val;  
        this.q = null;  
    }  
}
```

```
public void wait(Thread t) {  
    while(test&set(this.guard) == 1) {  
        // while busy do nothing  
    }  
    this.value -= 1;  
    if(this.value < 0) {  
        q.push(t);  
        t.sleep_and_reset_guard_to_0();  
    }  
    else {  
        this.guard = 0;  
    }  
}
```

```
public void signal() {  
    while(test&set(this.guard) == 1) {  
        // while busy do nothing  
    }  
    this.value += 1;  
    if(!q.isEmpty()) {// this.value <= 0  
        t = q.pop();  
        push_onto_ready_queue(t);  
    }  
    this.guard = 0;  
}
```

wait() and **signal()** are of course atomic!

either interrupts must be disabled or **test&set** used

Semaphore: Example

A

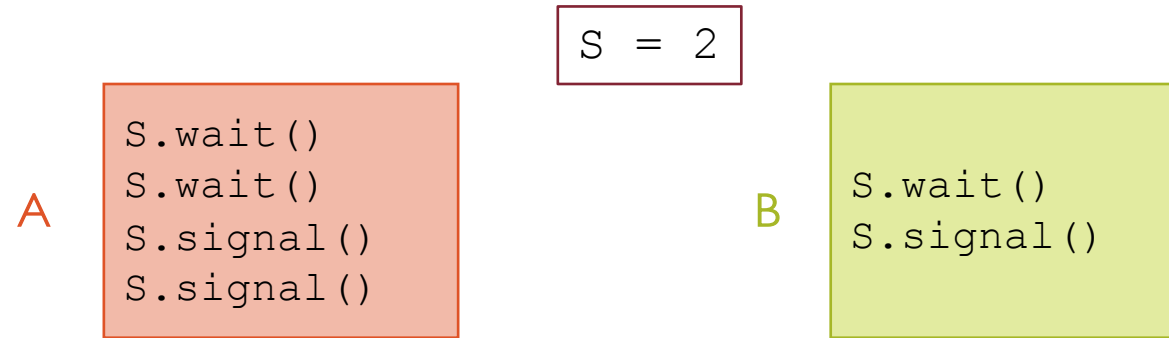
```
S.wait()  
S.wait()  
S.signal()  
S.signal()
```

S = 2

B

```
S.wait()  
S.signal()
```

Semaphore: Example



A possible execution flow

S (value)	Queue	A	B
2	∅	ready to exec	ready to exec

Semaphore: Example

A

```
S.wait()  
S.wait()  
S.signal()  
S.signal()
```

B

```
S.wait()  
S.signal()
```

A possible execution flow

A: S.wait()

S (value)	Queue	A	B
2	∅	ready to exec	ready to exec

Semaphore: Example

A

```
S.wait()  
S.wait()  
S.signal()  
S.signal()
```

B

```
S.wait()  
S.signal()
```

A possible execution flow

A: S.wait()

S (value)	Queue	A	B
2	∅	ready to exec	ready to exec
1	∅	ready to exec	ready to exec

Semaphore: Example

A

```
S.wait()  
S.wait()  
S.signal()  
S.signal()
```

B

```
S.wait()  
S.signal()
```

A possible execution flow

A: S.wait()

B: S.wait()

S (value)	Queue	A	B
2	∅	ready to exec	ready to exec
1	∅	ready to exec	ready to exec

Semaphore: Example

A

```
S.wait()  
S.wait()  
S.signal()  
S.signal()
```

B

```
S.wait()  
S.signal()
```

A possible execution flow

A: S.wait()

B: S.wait()

S (value)	Queue	A	B
2	∅	ready to exec	ready to exec
1	∅	ready to exec	ready to exec
0	∅	ready to exec	ready to exec

Semaphore: Example

A

```
S.wait()
S.wait()
S.signal()
S.signal()
```

B

```
S.wait()
S.signal()
```

A possible execution flow

A: S.wait()

B: S.wait()

A: S.wait()

S (value)	Queue	A	B
2	∅	ready to exec	ready to exec
1	∅	ready to exec	ready to exec
0	∅	ready to exec	ready to exec

Semaphore: Example

A

```
S.wait()
S.wait()
S.signal()
S.signal()
```

B

```
S.wait()
S.signal()
```

A possible execution flow

A: S.wait()

B: S.wait()

A: S.wait()

S (value)	Queue	A	B
2	∅	ready to exec	ready to exec
1	∅	ready to exec	ready to exec
0	∅	ready to exec	ready to exec
-1	A	blocked	ready to exec

Semaphore: Example

A

```
S.wait()
S.wait()
S.signal()
S.signal()
```

B

```
S.wait()
S.signal()
```

A possible execution flow

A: S.wait()
B: S.wait()
A: S.wait()
B: S.signal()

S (value)	Queue	A	B
2	∅	ready to exec	ready to exec
1	∅	ready to exec	ready to exec
0	∅	ready to exec	ready to exec
-1	A	blocked	ready to exec

Semaphore: Example

A

```
S.wait()
S.wait()
S.signal()
S.signal()
```

B

```
S.wait()
S.signal()
```

A possible execution flow

A: S.wait()

B: S.wait()

A: S.wait()

B: S.signal()

S (value)	Queue	A	B
2	∅	ready to exec	ready to exec
1	∅	ready to exec	ready to exec
0	∅	ready to exec	ready to exec
-1	A	blocked	ready to exec
0	∅	ready to exec	ready to exec

Semaphore: Example

A

```
S.wait()
S.wait()
S.signal()
S.signal()
```

B

```
S.wait()
S.signal()
```

A possible execution flow

	S (value)	Queue	A	B
	2	∅	ready to exec	ready to exec
A: S.wait()	1	∅	ready to exec	ready to exec
B: S.wait()	0	∅	ready to exec	ready to exec
A: S.wait()	-1	A	blocked	ready to exec
B: S.signal()	0	∅	ready to exec	ready to exec
A: S.signal()	1	∅	ready to exec	ready to exec
A: S.signal()	2	∅	ready to exec	ready to exec

Semaphores: Purposes

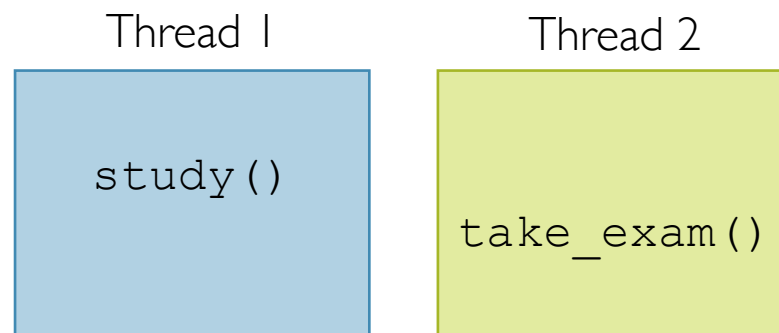
- **Mutual Exclusion:** used to guard critical sections
 - The initial value of the semaphore is set to 1
 - Call **wait()** before the critical section, **signal()** after the critical section

Semaphores: Purposes

- **Mutual Exclusion:** used to guard critical sections
 - The initial value of the semaphore is set to 1
 - Call **wait()** before the critical section, **signal()** after the critical section
- **Scheduling Constraints:** used to enforce threads to wait
 - The initial value of the semaphore is set to 0
 - Example → **join()** or **waitpid()**

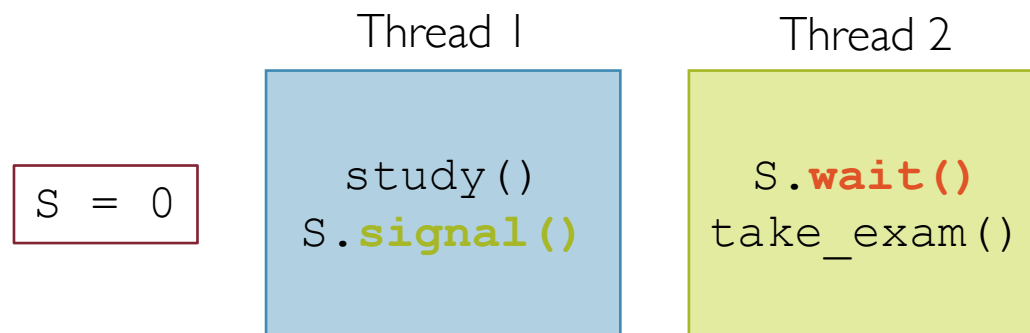
Semaphores: Purposes

- **Mutual Exclusion:** used to guard critical sections
 - The initial value of the semaphore is set to 1
 - Call **wait()** before the critical section, **signal()** after the critical section
- **Scheduling Constraints:** used to enforce threads to wait
 - The initial value of the semaphore is set to 0
 - Example → **join()** or **waitpid()**



Semaphores: Purposes

- **Mutual Exclusion:** used to guard critical sections
 - The initial value of the semaphore is set to 1
 - Call **wait()** before the critical section, **signal()** after the critical section
- **Scheduling Constraints:** used to enforce threads to wait
 - The initial value of the semaphore is set to 0
 - Example → **join()** or **waitpid()**



Producer-Consumer

Producer Process:

```
while (true)
{
    /* produce an item in nextProduced */
    while (counter == BUFFER_SIZE)
        ; /* do nothing */
    buffer[in] = nextProduced;
    in = (in + 1) % BUFFER_SIZE;
    counter++;
}
```

Consumer Process:

```
while (true)
{
    while (counter == 0)
        ; /* do nothing */
    nextConsumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    counter--;
    /* consume the item in nextConsumed */
}
```

Both the producer and the consumer share a **common buffer** (of items)

Producer-Consumer

Producer Process:

```
while (true)
{
    /* produce an item in nextProduced */
    while (counter == BUFFER_SIZE)
        ; /* do nothing */
    buffer[in] = nextProduced;
    in = (in + 1) % BUFFER_SIZE;
    counter++;
}
```

Consumer Process:

```
while (true)
{
    while (counter == 0)
        ; /* do nothing */
    nextConsumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    counter--;
    /* consume the item in nextConsumed */
}
```

Both the producer and the consumer share a **common buffer** (of items)

counter keeps track of the number of items currently in the buffer

Producer-Consumer

Producer Process:

```
while (true)
{
    /* produce an item in nextProduced */
    while (counter == BUFFER_SIZE)
        ; /* do nothing */
    buffer[in] = nextProduced;
    in = (in + 1) % BUFFER_SIZE;
    counter++;
}
```

Consumer Process:

```
while (true)
{
    while (counter == 0)
        ; /* do nothing */
    nextConsumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    counter--;
    /* consume the item in nextConsumed */
}
```

Both the producer and the consumer share a **common buffer** (of items)

counter keeps track of the number of items currently in the buffer

possible race condition as counter can be updated by the producer and consumer

Producer-Consumer: Race Condition

Producer:

```
register1 = counter  
register1 = register1 + 1  
counter = register1
```

Consumer:

```
register2 = counter  
register2 = register2 - 1  
counter = register2
```

Interleaving:

Assuming the initial value of counter is 5

T_0 :	producer	execute	$register_1 = counter$	{ $register_1 = 5$ }
T_1 :	producer	execute	$register_1 = register_1 + 1$	{ $register_1 = 6$ }
T_2 :	consumer	execute	$register_2 = counter$	{ $register_2 = 5$ }
T_3 :	consumer	execute	$register_2 = register_2 - 1$	{ $register_2 = 4$ }
T_4 :	producer	execute	$counter = register_1$	{ $counter = 6$ }
T_5 :	consumer	execute	$counter = register_2$	{ $counter = 4$ }

Producer-Consumer: Race Condition

Producer:

```
register1 = counter  
register1 = register1 + 1  
counter = register1
```

Consumer:

```
register2 = counter  
register2 = register2 - 1  
counter = register2
```

Interleaving:

Assuming the initial value of counter is 5

T ₀ :	producer	execute	register ₁ = counter	{register ₁ = 5}
T ₁ :	producer	execute	register ₁ = register ₁ + 1	{register ₁ = 6}
T ₂ :	consumer	execute	register ₂ = counter	{register ₂ = 5}
T ₃ :	consumer	execute	register ₂ = register ₂ - 1	{register ₂ = 4}
T ₄ :	producer	execute	counter = register ₁	{counter = 6}
T ₅ :	consumer	execute	counter = register ₂	{counter = 4}

Q1: What would be the resulting value of counter if the order of statements T4 and T5 were reversed?

Producer-Consumer: Race Condition

Producer:

```
register1 = counter  
register1 = register1 + 1  
counter = register1
```

Consumer:

```
register2 = counter  
register2 = register2 - 1  
counter = register2
```

Interleaving:

Assuming the initial value of counter is 5

T ₀ :	producer	execute	register ₁ = counter	{register ₁ = 5}
T ₁ :	producer	execute	register ₁ = register ₁ + 1	{register ₁ = 6}
T ₂ :	consumer	execute	register ₂ = counter	{register ₂ = 5}
T ₃ :	consumer	execute	register ₂ = register ₂ - 1	{register ₂ = 4}
T ₄ :	producer	execute	counter = register ₁	{counter = 6}
T ₅ :	consumer	execute	counter = register ₂	{counter = 4}

Q1: What would be the resulting value of counter if the order of statements T4 and T5 were reversed?

Q2: What *should* the value of counter be after one producer and one consumer, assuming the original value was 5?

Producer-Consumer: Desiderata

- Mutual Exclusion
 - Access to the shared buffer of items must be granted to a single thread at a time (either to the producer or the consumer)

Producer-Consumer: Desiderata

- Mutual Exclusion

- Access to the shared buffer of items must be granted to a single thread at a time (either to the producer or the consumer)

- Scheduling Constraints

- Producer can only produce a new item iff the buffer is **not full**
- Consumer can only consume an item iff the buffer is **not empty**

Producer-Consumer: Java Example

Semaphores: Wrap Up

- Generalization of locks
- Can be used for **3 purposes**:
 - To ensure mutually exclusive execution of a critical section as locks do (binary semaphore)
 - To control access to a shared pool of resources (counting semaphore)
 - To enforce scheduling constraints so as to execute threads according to some specific order

What's wrong with Semaphores?

- Not easy to get the meaning of waiting/signaling on a semaphore
- They are essentially shared global variables
- There is no direct connection between the semaphore and the data which the semaphore controls access to
- They serve multiple purposes (e.g., mutex, scheduling constraints, etc.)
- Their correctness depends on the programmer's ability

What's wrong with Semaphores?

- Not easy to get the meaning of waiting/signaling on a semaphore
- They are essentially shared global variables
- There is no direct connection between the semaphore and the data which the semaphore controls access to
- They serve multiple purposes (e.g., mutex, scheduling constraints, etc.)
- Their correctness depends on the programmer's ability

Solution: Use a higher level primitive called **monitors**

What is a Monitor?

- A monitor is a programming language construct that controls access to shared data

What is a Monitor?

- A monitor is a programming language construct that controls access to shared data
- Similar to a (Java/C++) class that embodies all together: **data**, **operations**, and **synchronization**

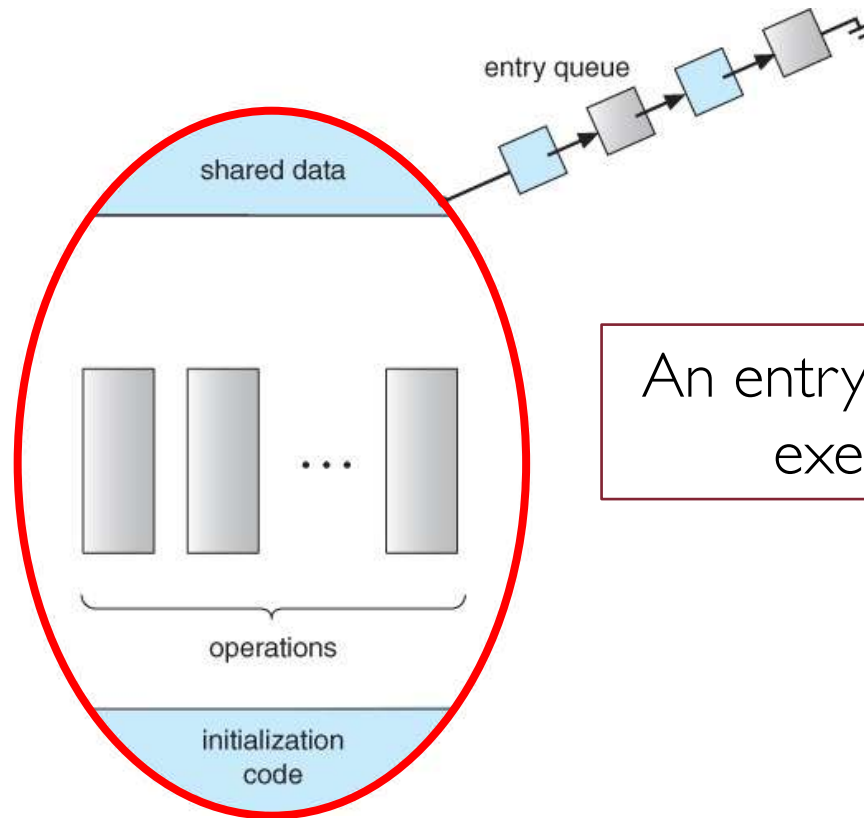
What is a Monitor?

- A monitor is a programming language construct that controls access to shared data
- Similar to a (Java/C++) class that embodies all together: **data**, **operations**, and **synchronization**
- Synchronization code added by compiler, enforced at runtime

What is a Monitor?

- A monitor is a programming language construct that controls access to shared data
- Similar to a (Java/C++) class that embodies all together: **data**, **operations**, and **synchronization**
- Synchronization code added by compiler, enforced at runtime
- Unlike classes, monitors:
 - guarantee mutual exclusion, i.e., only one thread may execute a monitor's method at a time
 - require all data to be private

Monitor: A Schematic Overview



An entry queue of processes waiting their turn to execute monitor operations (methods)

Monitor: A Formal Definition

- Defines a lock and zero or more **condition variables** for managing concurrent access to shared data

Monitor: A Formal Definition

- Defines a lock and zero or more **condition variables** for managing concurrent access to shared data
- Uses the lock to ensure that only a single thread is active within the monitor at any time

Monitor: A Formal Definition

- Defines a lock and zero or more **condition variables** for managing concurrent access to shared data
- Uses the lock to ensure that only a single thread is active within the monitor at any time
- The lock provides of course mutual exclusion for shared data

Monitor: Java Implementation Example

- It is straightforward to turn a Java class into a monitor by just:
 - Making all the data private
 - Making all methods (or non-private ones) **synchronized**

Monitor: Java Implementation Example

- It is straightforward to turn a Java class into a monitor by just:
 - Making all the data private
 - Making all methods (or non-private ones) **synchronized**
- The **synchronized** keyword indicates the method is subject to mutual exclusion

Monitor: Java Implementation Example

```
class Queue {  
    ...  
    private ArrayList<Item> data;  
    ...  
  
    public void synchronized add(Item i) {  
        data.add(i);  
    }  
  
    public Item synchronized remove() {  
        if (!data.isEmpty()) {  
            Item i = data.remove(0);  
            return i;  
        }  
    }  
}
```

Monitor: Java Implementation Example

```
class Queue {  
    ...  
    private ArrayList<Item> data;  
    ...  
  
    public void synchronized add(Item i) {  
        data.add(i);  
    }  
  
    public Item synchronized remove() {  
        if (!data.isEmpty()) {  
            Item i = data.remove(0);  
            return i;  
        }  
    }  
}
```

What happens if a thread tries to remove an element from an empty queue?

Condition Variables

- In the previous example, the **remove ()** method should wait until something is available on the queue

Condition Variables

- In the previous example, the **remove()** method should wait until something is available on the queue
 - Intuitively, the thread should sleep inside of the critical section

Condition Variables

- In the previous example, the **remove()** method should wait until something is available on the queue
 - Intuitively, the thread should sleep inside of the critical section
 - But if the thread sleeps while still holding a lock then no other threads can access the queue, add an item to it, and eventually wake up the sleeping thread

Condition Variables

- In the previous example, the **remove()** method should wait until something is available on the queue
 - Intuitively, the thread should sleep inside of the critical section
 - But if the thread sleeps while still holding a lock then no other threads can access the queue, add an item to it, and eventually wake up the sleeping thread
 - **Deadlock** (more on this later...)

Condition Variables

- In the previous example, the **remove()** method should wait until something is available on the queue
 - Intuitively, the thread should sleep inside of the critical section
 - But if the thread sleeps while still holding a lock then no other threads can access the queue, add an item to it, and eventually wake up the sleeping thread
 - **Deadlock** (more on this later...)
- Solution: **condition variables**
 - Enable a thread to sleep within a critical section
 - Any lock held by the thread is atomically released before going to sleep

Condition Variables: Operations

- Each condition variable supports 3 operations:

Condition Variables: Operations

- Each condition variable supports **3 operations**:
 - **wait** → release lock and go to sleep atomically (queue of waiters)

Condition Variables: Operations

- Each condition variable supports **3 operations**:
 - **wait** → release lock and go to sleep atomically (queue of waiters)
 - **signal** → wake up a waiting thread if one exists, otherwise it does nothing

Condition Variables: Operations

- Each condition variable supports **3 operations**:
 - **wait** → release lock and go to sleep atomically (queue of waiters)
 - **signal** → wake up a waiting thread if one exists, otherwise it does nothing
 - **broadcast** → wake up all waiting threads

Condition Variables: Operations

- Each condition variable supports **3 operations**:
 - **wait** → release lock and go to sleep atomically (queue of waiters)
 - **signal** → wake up a waiting thread if one exists, otherwise it does nothing
 - **broadcast** → wake up all waiting threads
- **Rule**: thread must hold the lock when doing condition variable operations
- **Note**: condition variables are not boolean objects!

Condition Variables in Java

- Use **wait()** to give up the lock

Condition Variables in Java

- Use **wait()** to give up the lock
- Use **notify()** to signal that the condition a thread is waiting on is satisfied

Condition Variables in Java

- Use **wait()** to give up the lock
- Use **notify()** to signal that the condition a thread is waiting on is satisfied
- Use **notifyAll()** to wake up all waiting threads

Condition Variables in Java

- Use **wait()** to give up the lock
- Use **notify()** to signal that the condition a thread is waiting on is satisfied
- Use **notifyAll()** to wake up all waiting threads
- Concretely, one condition variable per object

Monitor: Java Implementation Example

```
class Queue {  
    ...  
    private ArrayList<Item> data;  
    ...  
  
    public void synchronized add(Item i) {  
        data.add(i);  
        notify();  
    }  
  
    public Item synchronized remove() {  
        while (data.isEmpty()) {  
            wait(); // give up the lock and sleep  
        }  
        Item i = data.remove(0);  
        return i;  
    }  
}
```

Condition Variables are *not* Semaphores!

- Same operations yet entirely different semantics

Condition Variables are *not* Semaphores!

- Same operations yet entirely different semantics
- Access to the monitor is controlled by a lock

Condition Variables are *not* Semaphores!

- Same operations yet entirely different semantics
- Access to the monitor is controlled by a lock
- **wait()** blocks the calling thread, and gives up the lock
 - to call **wait()**, the thread has to be in the monitor (hence, it has the lock!)
 - on a semaphore, **wait()** just blocks the thread on the queue

Condition Variables are *not* Semaphores!

- Same operations yet entirely different semantics
- Access to the monitor is controlled by a lock
- **wait()** blocks the calling thread, and gives up the lock
 - to call **wait()**, the thread has to be in the monitor (hence, it has the lock!)
 - on a semaphore, **wait()** just blocks the thread on the queue
- **signal()** causes a waiting thread to wake up
 - If there is no waiting thread, the signal is lost though!
 - on a semaphore, signal increases the counter, allowing future entry even if no thread is currently waiting

`signal()`: Mesa- vs. Hoare-style Monitors

- **Mesa-style** (Nachos, Java, and most real OSs)
 - The signaling thread places a waiter on the ready queue, but signaler continues inside monitor
 - Condition is not necessarily true when waiter runs again
 - Returning from **`wait()`** is only a hint that something changed
 - Must re-check the conditional case

`signal()`: Mesa- vs. Hoare-style Monitors

- **Mesa-style** (Nachos, Java, and most real OSs)
 - The signaling thread places a waiter on the ready queue, but signaler continues inside monitor
 - Condition is not necessarily true when waiter runs again
 - Returning from **`wait()`** is only a hint that something changed
 - Must re-check the conditional case
- **Hoare-style** (most textbooks)
 - The signaling thread immediately switches to a waiting thread
 - The condition that the waiter was anticipating is guaranteed to hold when waiter

Mesa vs. Hoare Monitors

- Mesa-style

```
while (empty) {  
    wait(condition);  
}
```


Mesa vs. Hoare Monitors

- Mesa-style

```
while (empty) {  
    wait(condition);  
}
```

- Hoare-style

```
if (empty) {  
    wait(condition);  
}
```

Mesa vs. Hoare Monitors

- Mesa-style

```
while (empty) {  
    wait(condition);  
}
```

Easier to use and more efficient

- Hoare-style

```
if (empty) {  
    wait(condition);  
}
```

Easier to reason about the program's behaviour

Mesa vs. Hoare

Mesa

```
class Queue {  
    ...  
    private ArrayList<Item> data;  
    ...  
  
    public void synchronized add(Item i) {  
        data.add(i);  
        notify();  
    }  
  
    public Item synchronized remove() {  
        while (data.isEmpty()) {  
            wait(); // give up the lock and sleep  
        }  
        Item i = data.remove(0);  
        return i;  
    }  
}
```

The waiting thread may need to wait again after it is awakened, because some other thread could grab the lock and remove the item before it gets to run

Hoare

```
class Queue {  
    ...  
    private ArrayList<Item> data;  
    ...  
  
    public void synchronized add(Item i) {  
        data.add(i);  
        notify();  
    }  
  
    public Item synchronized remove() {  
        if (data.isEmpty()) {  
            wait(); // give up the lock and sleep  
        }  
        Item i = data.remove(0);  
        return i;  
    }  
}
```

The waiting thread runs immediately after an item is added to the queue

Readers-Writers Problem

- An object is shared among many threads, each belonging to one of two classes:
 - **Readers:** read data, never modify it
 - **Writers:** read data and modify it

Readers-Writers Problem

- An object is shared among many threads, each belonging to one of two classes:
 - **Readers:** read data, never modify it
 - **Writers:** read data and modify it
- Simplest solution:
 - Use a single lock on the data object for each operation
 - May be too restrictive!

Readers-Writers Problem

- An object is shared among many threads, each belonging to one of two classes:
 - **Readers:** read data, never modify it
 - **Writers:** read data and modify it
- Simplest solution:
 - Use a single lock on the data object for each operation
 - May be too restrictive!
- Each read or write of the shared data must happen within a critical section
- Guarantee mutual exclusion for writers
- Allow multiple readers to execute in the critical section at once

Readers-Writers Problem

- 2 variations of the problem depending on whether priority is on readers or writers

Readers-Writers Problem

- 2 variations of the problem depending on whether priority is on readers or writers
- first readers-writers problem (priority to readers)
 - if a reader wants access to the data, and there is not already a writer accessing it, then access is granted to the reader
 - possible starvation of the writers, as there could always be more readers coming along to access the data

Readers-Writers Problem

- **2 variations** of the problem depending on whether priority is on readers or writers
- **first readers-writers** problem (priority to readers)
 - if a reader wants access to the data, and there is not already a writer accessing it, then access is granted to the reader
 - possible starvation of the writers, as there could always be more readers coming along to access the data
- **second readers-writers** problem (priority to the writers)
 - when a writer wants access to the data it jumps to the head of the queue
 - possible starvation of the readers, as they are all blocked as long as there are writers

(First) Readers-Writers Problem: Solution I

- Use a counter and 2 binary semaphores
 - **numReaders** → used by the reader processes to count the number of readers currently accessing the data
 - **mutex** → binary semaphore used only by the readers for controlled access to **numReaders**
 - **rw_mutex** → binary semaphore used to block and release the writers

(First) Readers-Writers Problem: Solution I

- Use a counter and 2 binary semaphores
 - **numReaders** → used by the reader processes to count the number of readers currently accessing the data
 - **mutex** → binary semaphore used only by the readers for controlled access to **numReaders**
 - **rw_mutex** → binary semaphore used to block and release the writers
- The first reader to come along will block on **rw_mutex** if there is currently a writer accessing the data

(First) Readers-Writers Problem: Solution I

- Use a counter and 2 binary semaphores
 - **numReaders** → used by the reader processes to count the number of readers currently accessing the data
 - **mutex** → binary semaphore used only by the readers for controlled access to **numReaders**
 - **rw_mutex** → binary semaphore used to block and release the writers
- The first reader to come along will block on **rw_mutex** if there is currently a writer accessing the data
- All following readers will only block on **mutex** for their turn to increment **numReaders**

Solution I: Discussion

- The first reader blocks if there is a writer; any other reader who tries to enter gets blocked on **mutex**

Solution I: Discussion

- The first reader blocks if there is a writer; any other reader who tries to enter gets blocked on **mutex**
- Only the last reader to exit signals a waiting writer

Solution I: Discussion

- The first reader blocks if there is a writer; any other reader who tries to enter gets blocked on **mutex**
- Only the last reader to exit signals a waiting writer
- When a writer exits, if there is both a reader and writer waiting, which goes next depends on the scheduler

Solution I: Discussion

- The first reader blocks if there is a writer; any other reader who tries to enter gets blocked on **mutex**
- Only the last reader to exit signals a waiting writer
- When a writer exits, if there is both a reader and writer waiting, which goes next depends on the scheduler
- If a writer exits and a reader goes next, then all readers that are waiting will fall through (at least one is waiting on **rw_mutex** and zero or more can be waiting on **mutex**)

Solution I: Discussion

- The first reader blocks if there is a writer; any other reader who tries to enter gets blocked on **mutex**
- Only the last reader to exit signals a waiting writer
- When a writer exits, if there is both a reader and writer waiting, which goes next depends on the scheduler
- If a writer exits and a reader goes next, then all readers that are waiting will fall through (at least one is waiting on **rw_mutex** and zero or more can be waiting on **mutex**)
- Does this solution guarantee all threads will make progress?

Solution I: Discussion

- The first reader blocks if there is a writer; any other reader who tries to enter gets blocked on **mutex**
- Only the last reader to exit signals a waiting writer
- When a writer exits, if there is both a reader and writer waiting, which goes next depends on the scheduler
- If a writer exits and a reader goes next, then all readers that are waiting will fall through (at least one is waiting on **rw_mutex** and zero or more can be waiting on **mutex**)
- Does this solution guarantee all threads will make progress?
- Alternatively, let a writer enter its critical section first (priority to writers)

Solution II: Using Monitors

Summary

- 3 synchronization primitives:

Summary

- 3 synchronization primitives:
 - Locks → the simplest one

Summary

- 3 synchronization primitives:
 - Locks → the simplest one
 - Semaphores → a generalization of locks

Summary

- 3 synchronization primitives:
 - Locks → the simplest one
 - Semaphores → a generalization of locks
 - Monitors → highest-level primitives that wrap methods with mutex

Summary

- 3 synchronization primitives:
 - Locks → the simplest one
 - Semaphores → a generalization of locks
 - Monitors → highest-level primitives that wrap methods with mutex
- Examples of typical synchronization problems:

Summary

- 3 synchronization primitives:
 - Locks → the simplest one
 - Semaphores → a generalization of locks
 - Monitors → highest-level primitives that wrap methods with mutex
- Examples of typical synchronization problems:
 - Producers-Consumers

Summary

- 3 synchronization primitives:
 - Locks → the simplest one
 - Semaphores → a generalization of locks
 - Monitors → highest-level primitives that wrap methods with mutex
- Examples of typical synchronization problems:
 - Producers-Consumers
 - Readers-Writers