

# Systems and Networking I

Applied Computer Science and Artificial Intelligence  
2025-2026



**SAPIENZA**  
UNIVERSITÀ DI ROMA

Gabriele Tolomei

Computer Science Department

Sapienza Università di Roma

[tolomei@di.uniroma1.it](mailto:tolomei@di.uniroma1.it)

# Recap from Last Lecture

- Process is the **unit of execution** (running on a single CPU)
- OS keeps track of process-related information using an ad hoc data structure called **Process Control Block (PCB)**
- Process can be in one of **5 possible states**: **new**, **ready**, **waiting**, **running**, or **terminated**
- **Context switch** to intertwine the execution of multiple processes
- Process communication either via **message passing** or **shared memory**

# Today: CPU Scheduling

Policy to establish which process to execute on the CPU

- Basic scheduling concepts
- Scheduling **criteria/metrics**
- Scheduling **algorithms**
- Advanced scheduling concepts

# Basic Concepts

- Almost every program has some alternating cycles of CPU computations and I/O waiting
- Even a simple fetch from main memory takes a long time relatively to CPU speed
- Our assumptions: Multi-programmed, uni-processor system

# Basic Concepts

- In a system running a single process, the time spent waiting for I/O is wasted, and those CPU cycles are lost forever

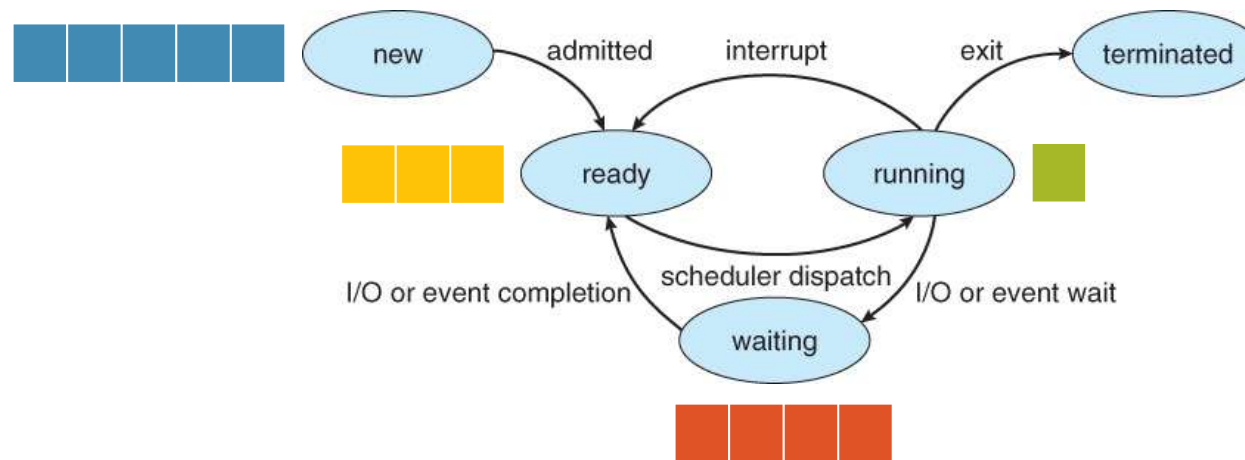
# Basic Concepts

- In a system running a single process, the time spent waiting for I/O is wasted, and those CPU cycles are lost forever
- A scheduling system allows one process to use the CPU while another is waiting for I/O, thereby maximizing system utilization

# Basic Concepts

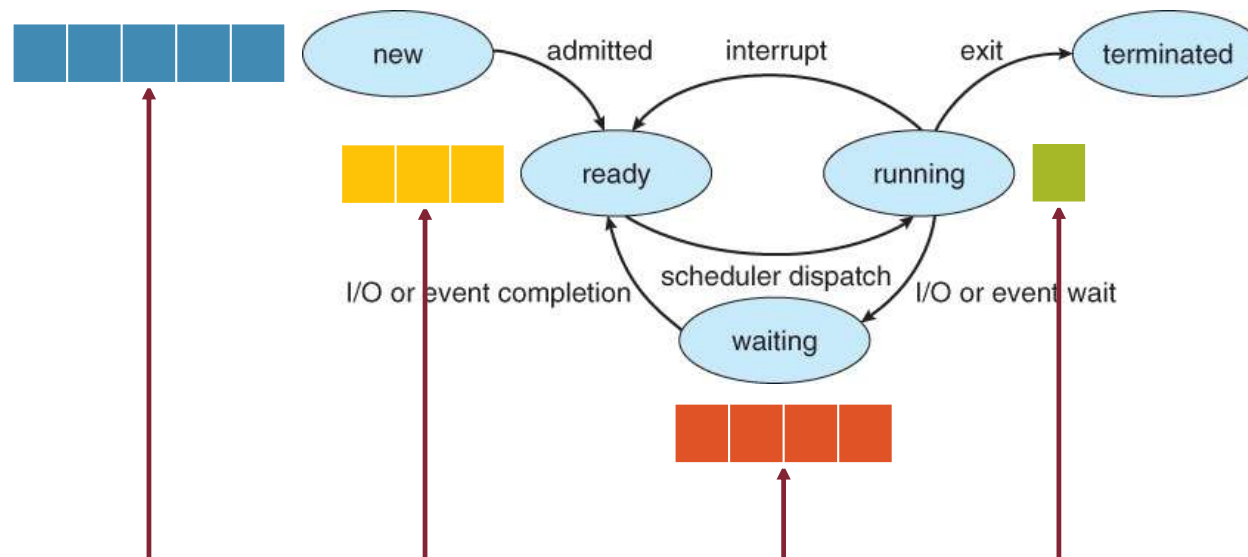
- In a system running a single process, the time spent waiting for I/O is wasted, and those CPU cycles are lost forever
- A scheduling system allows one process to use the CPU while another is waiting for I/O, thereby maximizing system utilization
- **Challenge:** Make the system as "efficient" and "fair" as possible, subject to varying and often dynamic conditions

# Process Execution State Diagram



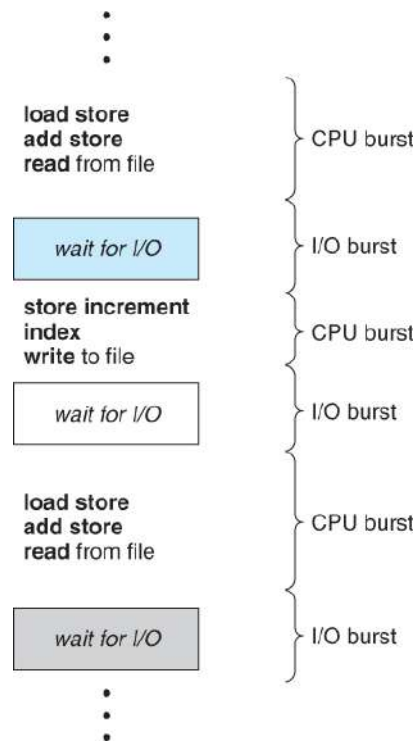


# Process Execution State Diagram



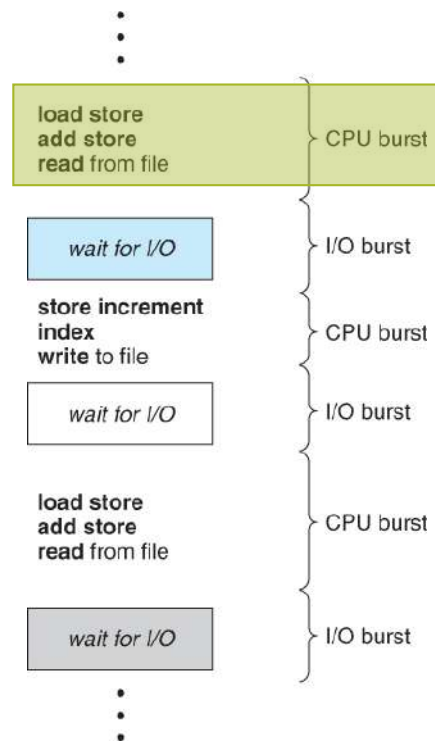
Processes managed by the OS reside in exactly one of the state queues

# CPU vs. I/O Burst Cycle



All processes alternate between two states in a continuing cycle: CPU burst and I/O burst

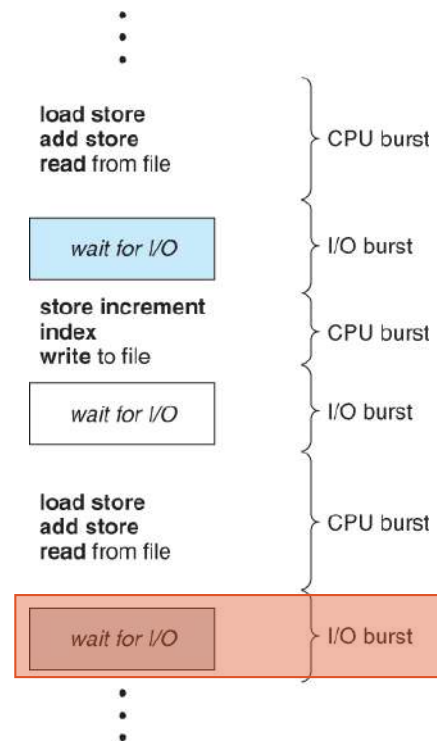
# CPU vs. I/O Burst Cycle



All processes alternate between two states in a continuing cycle: **CPU burst** and **I/O burst**

**CPU burst** → performing calculations

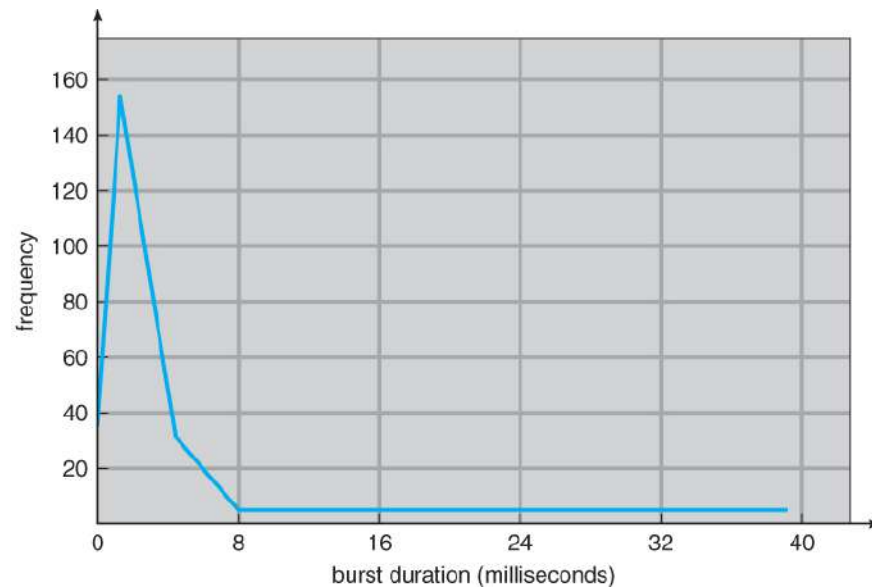
# CPU vs. I/O Burst Cycle



All processes alternate between two states in a continuing cycle: CPU burst and I/O burst

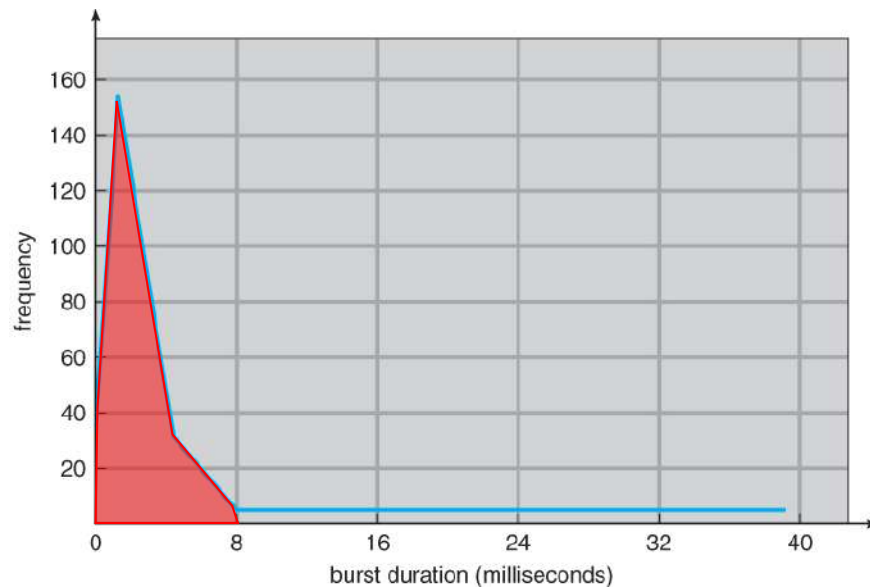
I/O burst → waiting for data transfer in or out of the system

# CPU Burst Cycle: Frequency Pattern



Highly skewed distribution

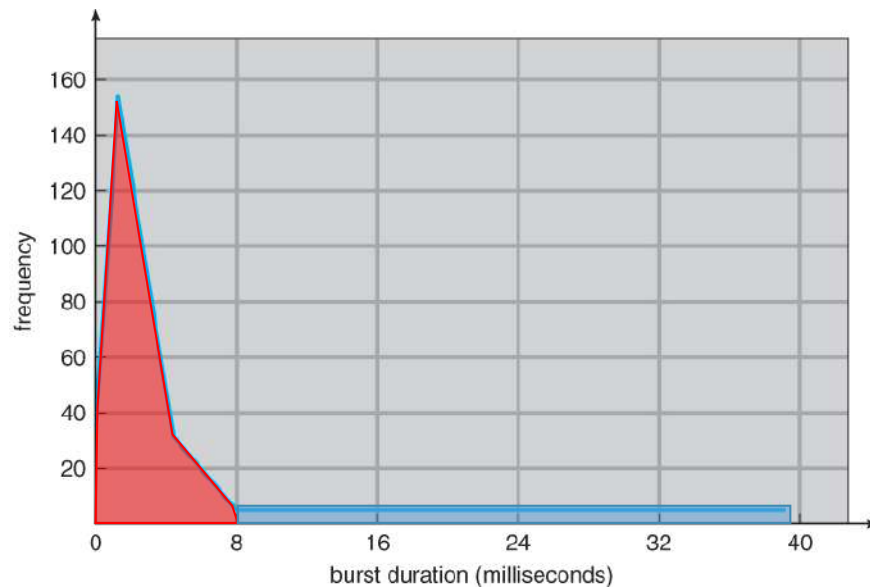
# CPU Burst Cycle: Frequency Pattern



Highly skewed distribution

The vast majority of processes have **short** CPU bursts

# CPU Burst Cycle: Frequency Pattern



Highly skewed distribution

The vast majority of processes have **short** CPU bursts

Few processes exhibit **very long** CPU bursts

# Long- vs. Short-term Scheduling

## Long-term scheduling

How does the OS determine the level of multiprogramming (i.e., the number of processes to be loaded in main memory)



# Long- vs. Short-term Scheduling

## Long-term scheduling

How does the OS determine the level of multiprogramming (i.e., the number of processes to be loaded in main memory)

## Short-term scheduling

How does the OS select a process to execute from the ready queue

# Long- vs. Short-term Scheduling

## Long-term scheduling

How does the OS determine the level of multiprogramming (i.e., the number of processes to be loaded in main memory)

## Short-term scheduling

How does the OS select a process to execute from the ready queue

CPU Scheduler

# CPU Scheduler

- Whenever the CPU becomes idle, it picks another process from the **ready** queue to execute next
- The data structure for the ready queue and the algorithm used to select the next process are not necessarily based on a FIFO queue

# CPU Scheduler

- Whenever the CPU becomes idle, it picks another process from the **ready** queue to execute next
- The data structure for the ready queue and the algorithm used to select the next process are not necessarily based on a FIFO queue

Policy goals vs. Mechanism implementations

# CPU Scheduling: When?

CPU scheduling decisions take place under one of 4 conditions:

# CPU Scheduling: When?

CPU scheduling decisions take place under one of 4 conditions:

1. When a process switches from the **running** state to the **waiting** state

# CPU Scheduling: When?

CPU scheduling decisions take place under one of 4 conditions:

1. When a process switches from the **running** state to the **waiting** state

for an I/O request or  
invocation of the wait  
system call

# CPU Scheduling: When?

CPU scheduling decisions take place under one of 4 conditions:

1. When a process switches from the running state to the waiting state
2. When a process switches from the **running** state to the **ready** state



# CPU Scheduling: When?

CPU scheduling decisions take place under one of 4 conditions:

1. When a process switches from the running state to the waiting state
2. When a process switches from the **running** state to the **ready** state

in response to an interrupt

# CPU Scheduling: When?

CPU scheduling decisions take place under one of 4 conditions:

1. When a process switches from the running state to the waiting state
2. When a process switches from the running state to the ready state
3. When a process switches from the **waiting** state to the **ready** state

# CPU Scheduling: When?

CPU scheduling decisions take place under one of 4 conditions:

1. When a process switches from the running state to the waiting state
2. When a process switches from the running state to the ready state
3. When a process switches from the **waiting** state to the **ready** state

after I/O completion  
or a return from wait

# CPU Scheduling: When?

CPU scheduling decisions take place under one of 4 conditions:

1. When a process switches from the running state to the waiting state
2. When a process switches from the running state to the ready state
3. When a process switches from the waiting state to the ready state
4. When a process is **created** or **terminates**

# CPU Scheduling: When?

CPU scheduling decisions take place under one of 4 conditions:

1. When a process switches from the running state to the waiting state
2. When a process switches from the running state to the ready state
3. When a process switches from the waiting state to the ready state
4. When a process is created or terminates

No choice!  
A new process **must** be selected

# CPU Scheduling: When?

CPU scheduling decisions take place under one of 4 conditions:

1. When a process switches from the running state to the waiting state
2. When a process switches from the running state to the ready state
3. When a process switches from the waiting state to the ready state
4. When a process is created or terminates

Either continue with the current process or select a new one

# Non-preemptive vs. Preemptive

## Non-preemptive scheduling

If it takes place only when there is no choice (i.e., conditions 1 and 4)  
Once a process starts it keeps running until it either voluntarily blocks or it finishes

# Non-preemptive vs. Preemptive

## Non-preemptive scheduling

If it takes place only when there is no choice (i.e., conditions 1 and 4)  
Once a process starts it keeps running until it either voluntarily blocks or it finishes

## Preemptive scheduling

Whenever scheduling takes also place under conditions 2 and 3



# Non-preemptive vs. Preemptive: Examples

	Windows	Mac	UNIX-like
Non-preemptive	up to Win 3.x	up to Mac OS 9.x	-
Preemptive	since Win 95	since Mac OS X	since forever

# Preemption: Issues

- Preemption might cause troubles if it occurs while:
  - the kernel is busy implementing a system call (e.g., updating critical kernel data structures)
  - two processes share data, one may get interrupted in the middle of updating shared data structures

# Preemption: Issues

- Possible countermeasures:
  - Make the process wait until the system call has either completed or blocked before allowing the preemption

# Preemption: Issues

- Possible countermeasures:
  - Make the process wait until the system call has either completed or blocked before allowing the preemption

problematic for real-time systems, as real-time response can no longer be guaranteed

# Preemption: Issues

- Possible countermeasures:
  - Make the process wait until the system call has either completed or blocked before allowing the preemption

problematic for real-time systems, as real-time response can no longer be guaranteed

- Disable interrupts before entering critical code section and re-enabling immediately afterwards

# Preemption: Issues

- Possible countermeasures:
  - Make the process wait until the system call has either completed or blocked before allowing the preemption

problematic for real-time systems, as real-time response can no longer be guaranteed

- Disable interrupts before entering critical code section and re-enabling immediately afterwards

should only be done in rare situations, and only on very short pieces of code that will finish quickly

# The Dispatcher

- The module that gives control of the CPU to the process selected by the scheduler

# The Dispatcher

- The module that gives control of the CPU to the process selected by the scheduler
- Its functions include:
  - Switching context
  - Switching to user mode
  - Jumping to the proper location in the newly loaded program



# The Dispatcher

- The module that gives control of the CPU to the process selected by the scheduler
- Its functions include:
  - Switching context
  - Switching to user mode
  - Jumping to the proper location in the newly loaded program
- The dispatcher is run on every context switch therefore the time it consumes (**dispatch latency**) must be as shortest as possible

# Useful Definitions

- **Arrival Time:** time at which the process arrives in the **ready** queue

# Useful Definitions

- **Arrival Time:** time at which the process arrives in the ready queue
- **Completion Time:** time at which the process completes its execution

# Useful Definitions

- **Arrival Time:** time at which the process arrives in the ready queue
- **Completion Time:** time at which the process completes its execution
- **Burst Time:** time required by a process for **CPU execution**

# Useful Definitions

- **Arrival Time:** time at which the process arrives in the ready queue
- **Completion Time:** time at which the process completes its execution
- **Burst Time:** time required by a process for CPU execution
- **Turnaround Time:** time difference between completion and arrival time

# Useful Definitions

- **Arrival Time:** time at which the process arrives in the ready queue
- **Completion Time:** time at which the process completes its execution
- **Burst Time:** time required by a process for CPU execution
- **Turnaround Time:** time difference between completion and arrival time
- **Waiting Time:** time difference between turnaround time and burst time

# Useful Definitions

- **Arrival Time:** time at which the process arrives in the ready queue
- **Completion Time:** time at which the process completes its execution
- **Burst Time:** time required by a process for CPU execution
- **Turnaround Time:** time difference between completion and arrival time
- **Waiting Time:** time difference between turnaround time and burst time

## NOTE

I/O waiting time is **not** considered here!

# Useful Definitions

$T^{arrival}$  = arrival time

$T^{completion}$  = completion time

$T^{burst}$  = burst time

$T^{turnaround}$  = turnaround time =  $T^{completion} - T^{arrival}$

$T^{waiting}$  = waiting time =  $T^{turnaround} - T^{burst}$



# Scheduling Criteria/Metrics

- There are several different criteria to consider when trying to select the "best" scheduling algorithm, including:
  - CPU utilization
  - Throughput
  - Tournaround time
  - Waiting time
  - Response time

# Scheduling Criteria/Metrics

- There are several different criteria to consider when trying to select the "best" scheduling algorithm, including:

Intuitively, the percentage of time the CPU is busy

- **CPU utilization**
- Throughput
- Turnaround time
- Waiting time
- Response time

# Scheduling Criteria/Metrics

- There are several different criteria to consider when trying to select the "best" scheduling algorithm, including:

- **CPU utilization**

- Throughput
  - Turnaround time
  - Waiting time
  - Response time

Intuitively, the percentage of time the CPU is busy

Ideally the CPU would be busy 100% of the time, so as to waste 0 CPU cycles

# Scheduling Criteria/Metrics

- There are several different criteria to consider when trying to select the "best" scheduling algorithm, including:

- **CPU utilization**

Intuitively, the percentage of time the CPU is busy

- Throughput

Ideally the CPU would be busy 100% of the time, so as to waste 0 CPU cycles

- Turnaround time

- Waiting time

On a real system CPU usage should range from 40% (lightly loaded) to 90% (heavily loaded)

- Response time

# Scheduling Criteria/Metrics

- There are several different criteria to consider when trying to select the "best" scheduling algorithm, including:

- **CPU utilization**

Intuitively, the percentage of time the CPU is busy

- Throughput

Ideally the CPU would be busy 100% of the time, so as to waste 0 CPU cycles

- Turnaround time

- Waiting time

On a real system CPU usage should range from 40% (lightly loaded) to 90% (heavily loaded)

- Response time

**maximize**

# Scheduling Criteria/Metrics

- There are several different criteria to consider when trying to select the "best" scheduling algorithm, including:

- CPU utilization
- **Throughput**
- Turnaround time
- Waiting time
- Response time

The number of processes completed in a unit of time

# Scheduling Criteria/Metrics

- There are several different criteria to consider when trying to select the "best" scheduling algorithm, including:

- CPU utilization
- **Throughput**
- Turnaround time
- Waiting time
- Response time

The number of processes  
completed in a unit of time

May range from 10 per second  
to 1 per hour

# Scheduling Criteria/Metrics

- There are several different criteria to consider when trying to select the "best" scheduling algorithm, including:

- CPU utilization
- **Throughput**
- Turnaround time
- Waiting time
- Response time

The number of processes  
completed in a unit of time

May range from 10 per second  
to 1 per hour

**maximize**



# Scheduling Criteria/Metrics

- There are several different criteria to consider when trying to select the "best" scheduling algorithm, including:

- CPU utilization
- Throughput
- **Tournaround time**
- Waiting time
- Response time

Time required for a particular process to complete, from submission time to completion

# Scheduling Criteria/Metrics

- There are several different criteria to consider when trying to select the "best" scheduling algorithm, including:

- CPU utilization
- Throughput
- **Tournaround time**
- Waiting time
- Response time

Time required for a particular process to complete, from submission time to completion

Includes all the waiting time

# Scheduling Criteria/Metrics

- There are several different criteria to consider when trying to select the "best" scheduling algorithm, including:

- CPU utilization
- Throughput
- **Tournaround time**
- Waiting time
- Response time

Time required for a particular process to complete, from submission time to completion

Includes all the waiting time

minimize

# Scheduling Criteria/Metrics

- There are several different criteria to consider when trying to select the "best" scheduling algorithm, including:

- CPU utilization
- Throughput
- Turnaround time
- **Waiting time**
- Response time

How much time processes spend in the ready queue waiting their turn to get on the CPU

# Scheduling Criteria/Metrics

- There are several different criteria to consider when trying to select the "best" scheduling algorithm, including:

- CPU utilization
- Throughput
- Turnaround time
- **Waiting time**
- Response time

How much time processes spend in the ready queue waiting their turn to get on the CPU

**Not** to be confused with the waiting state!

# Scheduling Criteria/Metrics

- There are several different criteria to consider when trying to select the "best" scheduling algorithm, including:

- CPU utilization
- Throughput
- Turnaround time
- **Waiting time**
- Response time

How much time processes spend in the ready queue waiting their turn to get on the CPU

**Not** to be confused with the waiting state!

Processes in the waiting state are not under control of the CPU scheduler (they are just not ready to run!)

# Scheduling Criteria/Metrics

- There are several different criteria to consider when trying to select the "best" scheduling algorithm, including:

- CPU utilization
- Throughput
- Turnaround time
- **Waiting time**
- Response time

How much time processes spend in the ready queue waiting their turn to get on the CPU

**Not** to be confused with the waiting state!

Processes in the waiting state are not under control of the CPU scheduler (they are just not ready to run!)

minimize

# Scheduling Criteria/Metrics

- There are several different criteria to consider when trying to select the "best" scheduling algorithm, including:

- CPU utilization
- Throughput
- Turnaround time
- Waiting time
- **Response time**

The time taken from the issuance of a command to the beginning of a response to that command



# Scheduling Criteria/Metrics

- There are several different criteria to consider when trying to select the "best" scheduling algorithm, including:

- CPU utilization
- Throughput
- Turnaround time
- Waiting time
- **Response time**

The time taken from the issuance of a command to the beginning of a response to that command

Mostly, for interactive processes

# Scheduling Criteria/Metrics

- There are several different criteria to consider when trying to select the "best" scheduling algorithm, including:

- CPU utilization
- Throughput
- Turnaround time
- Waiting time
- **Response time**

The time taken from the issuance of a command to the beginning of a response to that command

Mostly, for interactive processes

minimize

# Scheduling: Trade-off

- Ideally, choose a CPU scheduler that optimizes all metrics simultaneously
- Generally, the above is impossible and a trade-off is needed!
- **Idea:** Choose a scheduling algorithm based on its ability to satisfy a given policy

# Scheduling Policies

- Minimize **average** response time
  - Provide output to the user as quickly as possible

# Scheduling Policies

- Minimize **average** response time
  - Provide output to the user as quickly as possible
- Minimize the **maximum** response time
  - Worst-case bound

# Scheduling Policies

- Minimize **average** response time
  - Provide output to the user as quickly as possible
- Minimize the **maximum** response time
  - Worst-case bound
- Minimize **variance** of response time
  - Users are more accepting of a consistent predictable system than an inconsistent one

# Scheduling Policies

- Minimize **average** response time
  - Provide output to the user as quickly as possible
- Minimize the **maximum** response time
  - Worst-case bound
- Minimize **variance** of response time
  - Users are more accepting of a consistent predictable system than an inconsistent one

Typical of **interactive systems**

# Scheduling Policies

- Maximize throughput translates into:
  - Minimizing overhead (OS context switching)
  - Efficient usage of system resources (CPU and I/O devices)



# Scheduling Policies

- Maximize throughput translates into:
  - Minimizing overhead (OS context switching)
  - Efficient usage of system resources (CPU and I/O devices)
- Minimize waiting time
  - Give every process the same amount of time on the CPU
  - Might increase the average response time

# Scheduling Policies

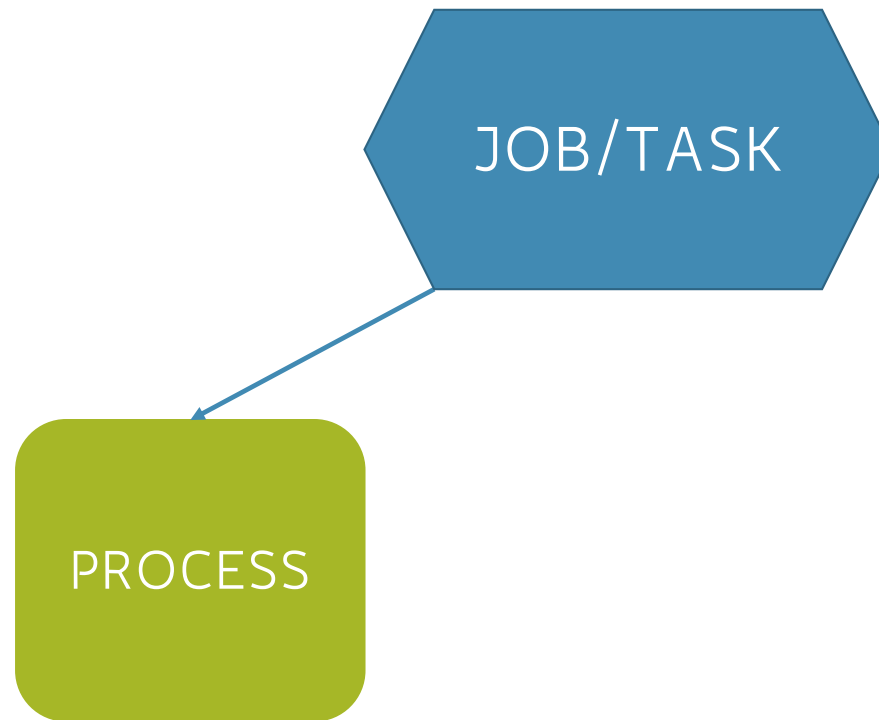
- Maximize throughput translates into:
  - Minimizing overhead (OS context switching)
  - Efficient usage of system resources (CPU and I/O devices)
- Minimize waiting time
  - Give every process the same amount of time on the CPU
  - Might increase the average response time

Typical of batch systems


# A Quick Note on Terminology



# A Quick Note on Terminology



# A Quick Note on Terminology



# Scheduling Policies: Our Assumptions

- One process per user

# Scheduling Policies: Our Assumptions

- One process per user
- Processes are independent from each other (i.e., no communication)

# Scheduling Policies: Our Assumptions

- One process per user
- Processes are independent from each other (i.e., no communication)
- One single thread (of execution) per process



# Scheduling Policies: Our Assumptions

- One process per user
- Processes are independent from each other (i.e., no communication)
- One single thread (of execution) per process

We will talk about **threads** very soon but for now most of the things we will be discussing remain valid even on a multi-threaded system

# Summary

- Scheduling allows one process to use the CPU while another is waiting for I/O

# Summary

- Scheduling allows one process to use the CPU while another is waiting for I/O
- The goal of scheduling is to maximize system utilization

# Summary

- Scheduling allows one process to use the CPU while another is waiting for I/O
- The goal of scheduling is to maximize system utilization
- **non-preemptive** vs. **preemptive** scheduler

# Summary

- Scheduling allows one process to use the CPU while another is waiting for I/O
- The goal of scheduling is to maximize system utilization
- **non-preemptive** vs. **preemptive** scheduler
- Different scheduling policies optimize different metrics