

# Systems and Networking I

Applied Computer Science and Artificial Intelligence

2025-2026



SAPIENZA  
UNIVERSITÀ DI ROMA

Gabriele Tolomei

Computer Science Department  
Sapienza Università di Roma

[tolomei@di.uniroma1.it](mailto:tolomei@di.uniroma1.it)

# Recap from Last Lecture

- Synchronization primitives:
  - Locks
  - Semaphores
  - Monitors

# Recap from Last Lecture

- Synchronization primitives:
  - Locks
  - Semaphores
  - Monitors
- 2 fundamental synchronization problems:
  - Producers-Consumers
  - Readers-Writers

# Our Journey

- What is deadlock?
- Conditions for deadlock to happen
- Deadlock prevention
- Deadlock avoidance
- Deadlock detection and recovery

# Our Journey

- What is deadlock?
- Conditions for deadlock to happen
- Deadlock prevention
- Deadlock avoidance
- Deadlock detection and recovery

# What is Deadlock?

"When two trains approach each other at a crossing,  
both shall come to a full stop and neither shall start  
up again until the other has gone."

Kansas legislation early 1900's

# What is Deadlock?

Intuitively, a condition where two or more threads are waiting for an event that can only be generated by the very same threads

# What is Deadlock?

Intuitively, a condition where two or more threads are waiting for an event that can only be generated by the very same threads

Thread A

```
printer.wait();
disk.wait();

// copy from disk to printer

printer.signal();
disk.signal();
```

Thread B

```
disk.wait();
printer.wait();

// copy from disk to printer

printer.signal();
disk.signal();
```

# What is Deadlock?

Intuitively, a condition where two or more threads are waiting for an event that can only be generated by the very same threads

Thread A

```
printer.wait();
disk.wait();

// copy from disk to printer

printer.signal();
disk.signal();
```

A starts first

Thread B

```
disk.wait();
printer.wait();

// copy from disk to printer

printer.signal();
disk.signal();
```

# What is Deadlock?

Intuitively, a condition where two or more threads are waiting for an event that can only be generated by the very same threads

Thread A

```
printer.wait(); Acquires printer and context switch  
disk.wait();  
  
// copy from disk to printer  
  
printer.signal();  
disk.signal();
```

Thread B

```
disk.wait();  
printer.wait();  
  
// copy from disk to printer  
  
printer.signal();  
disk.signal();
```

# What is Deadlock?

Intuitively, a condition where two or more threads are waiting for an event that can only be generated by the very same threads

Thread A

```
printer.wait();
disk.wait();

// copy from disk to printer

printer.signal();
disk.signal();
```

Thread B

B takes over

```
disk.wait();
printer.wait();

// copy from disk to printer

printer.signal();
disk.signal();
```

# What is Deadlock?

Intuitively, a condition where two or more threads are waiting for an event that can only be generated by the very same threads

Thread A

```
printer.wait();
disk.wait();

// copy from disk to printer

printer.signal();
disk.signal();
```

Thread B

```
disk.wait();           Acquires disk and
                     context switch
                     ↓
printer.wait();

// copy from disk to printer

printer.signal();
disk.signal();
```

# What is Deadlock?

Intuitively, a condition where two or more threads are waiting for an event that can only be generated by the very same threads

Thread A

```
printer.wait();
disk.wait();

// copy from disk to printer

printer.signal();
disk.signal();
```

A executes again and blocks

Thread B

```
disk.wait();
printer.wait();

// copy from disk to printer

printer.signal();
disk.signal();
```

# What is Deadlock?

Intuitively, a condition where two or more threads are waiting for an event that can only be generated by the very same threads

Thread A

```
printer.wait();
disk.wait();

// copy from disk to printer

printer.signal();
disk.signal();
```

Thread B

```
disk.wait();      B executes again and
printer.wait();    blocks

// copy from disk to printer

printer.signal();
disk.signal();
```

# What is Deadlock?

Intuitively, a condition where two or more threads are waiting for an event that can only be generated by the very same threads

Thread A

```
printer.wait();
disk.wait();

// copy from disk to printer

printer.signal();
disk.signal();
```

Thread B

```
disk.wait();
printer.wait();

// copy from disk to printer

printer.signal();
disk.signal();
```

A waits B to release the disk

# What is Deadlock?

Intuitively, a condition where two or more threads are waiting for an event that can only be generated by the very same threads

Thread A

```
printer.wait();
disk.wait();

// copy from disk to printer

printer.signal();
disk.signal();
```

Thread B

```
disk.wait();
printer.wait();

// copy from disk to printer

printer.signal();
disk.signal();
```

B waits A to release the printer

# Deadlock: Terminology

- **Deadlock:** it can occur when multiple threads compete for a finite number of resources

# Deadlock: Terminology

- **Deadlock:** it can occur when multiple threads compete for a finite number of resources
- **Deadlock prevention (offline):** imposes restrictions/rules on how to write deadlock-free programs

# Deadlock: Terminology

- **Deadlock:** it can occur when multiple threads compete for a finite number of resources
- **Deadlock prevention (offline):** imposes restrictions/rules on how to write deadlock-free programs
- **Deadlock avoidance (online):** scheduling threads to avoid deadlocks

# Deadlock: Terminology

- **Deadlock:** it can occur when multiple threads compete for a finite number of resources
- **Deadlock prevention (offline):** imposes restrictions/rules on how to write deadlock-free programs
- **Deadlock avoidance (online):** scheduling threads to avoid deadlocks
- **Deadlock detection (online):** finds instances of deadlocks and tries to recover

# Deadlock vs. Starvation

- Not to be confused with each other!

# Deadlock vs. Starvation

- Not to be confused with each other!
- Related terms but each one refers to a specific situation

# Deadlock vs. Starvation

- Not to be confused with each other!
- Related terms but each one refers to a specific situation
- Starvation occurs when a thread waits indefinitely for some resource but other threads are actually making progress using that resource

# Deadlock vs. Starvation

- Not to be confused with each other!
- Related terms but each one refers to a specific situation
- Starvation occurs when a thread waits indefinitely for some resource but other threads are actually making progress using that resource
- The main difference with deadlock is that the system is not completely stuck!

# Our Journey

- What is deadlock?
- Conditions for deadlock to happen
- Deadlock prevention
- Deadlock avoidance
- Deadlock detection and recovery

# Necessary Conditions for Deadlock

- Deadlock can happen if **all the 4 conditions** below hold

# Necessary Conditions for Deadlock

- Deadlock can happen if **all the 4 conditions** below hold
  - Mutual Exclusion → at least one thread must hold a non-sharable resource (e.g., only one thread grabs a lock)

# Necessary Conditions for Deadlock

- Deadlock can happen if **all the 4 conditions** below hold
  - **Mutual Exclusion** → at least one thread must hold a non-sharable resource (e.g., only one thread grabs a lock)
  - **Hold-and-Wait** → at least one thread is holding a non-sharable resource (e.g., a lock) and is waiting for other resource(s) to become available (e.g., other locks to acquire)

# Necessary Conditions for Deadlock

- Deadlock can happen if **all the 4 conditions** below hold
  - **Mutual Exclusion** → at least one thread must hold a non-sharable resource (e.g., only one thread grabs a lock)
  - **Hold-and-Wait** → at least one thread is holding a non-sharable resource (e.g., a lock) and is waiting for other resource(s) to become available (e.g., other locks to acquire)
  - **No Preemption** → a thread can only release a resource voluntarily; neither another thread nor the OS can force it to release the resource

# Necessary Conditions for Deadlock

- Deadlock can happen if **all the 4 conditions** below hold
  - **Mutual Exclusion** → at least one thread must hold a non-sharable resource (e.g., only one thread grabs a lock)
  - **Hold-and-Wait** → at least one thread is holding a non-sharable resource (e.g., a lock) and is waiting for other resource(s) to become available (e.g., other locks to acquire)
  - **No Preemption** → a thread can only release a resource voluntarily; neither another thread nor the OS can force it to release the resource
  - **Circular Wait** → a circular chain of waiting threads  $t_1, \dots, t_n$  where  $t_i$  holds a resource requested by  $t_{(i+1)\%n}$

# Our Journey

- What is deadlock?
- Conditions for deadlock to happen
- **Deadlock prevention**
- Deadlock avoidance
- Deadlock detection and recovery

# Deadlock Prevention

- Ensure that **at least one** of the **4** necessary conditions doesn't hold
  - Mutual Exclusion → make all resources sharable
    - Not all can be shared
    - E.g., disks, printers, etc.

# Deadlock Prevention

- Ensure that **at least one** of the **4** necessary conditions doesn't hold
  - Hold-and-Wait → a thread cannot hold a lock when it requests another
    - Acquire all locks at once, atomically
    - Use a global lock that wraps the acquisition of all locks
    - Hard to predict all the resources a thread will need and inefficient!

# Deadlock Prevention

- Ensure that **at least one** of the **4** necessary conditions doesn't hold
  - No Preemption → if a thread requests a resource that cannot be allocated to it, the OS preempts (releases) all the resources that the thread is already holding
    - Some thread libraries allow "trying" acquiring multiple locks
    - Not all resources can be easily preempted (e.g., printers)

# Deadlock Prevention

- Ensure that **at least one** of the **4** necessary conditions doesn't hold
  - Circular Wait → impose an ordering (i.e., numbering) on resources and enforce to request them in such order
    - Hard to establish such an order

# Our Journey

- What is deadlock?
- Conditions for deadlock to happen
- Deadlock prevention
- **Deadlock avoidance**
- Deadlock detection and recovery

# Deadlock Avoidance via Scheduling

- An alternative to statically preventing deadlock upfront

# Deadlock Avoidance via Scheduling

- An alternative to statically preventing deadlock upfront
- Avoidance requires some global knowledge of which locks various threads can grab

# Deadlock Avoidance via Scheduling

- An alternative to statically preventing deadlock upfront
- Avoidance requires some global knowledge of which locks various threads can grab
- Based on that knowledge, the OS will schedule threads to guarantee that no deadlock occurs

# Deadlock Avoidance via Scheduling

- An alternative to statically preventing deadlock upfront
- Avoidance requires some global knowledge of which locks various threads can grab
- Based on that knowledge, the OS will schedule threads to guarantee that no deadlock occurs
- Can be used only in limited environments where one has full knowledge of all tasks and locks needed

# Deadlock Avoidance: Example

- 4 threads:  $T_1, T_2, T_3, T_4$
- 2 CPUs:  $CPU_1, CPU_2$
- knowledge:
  - $T_1$  grabs locks  $L_1$  and  $L_2$  (in some order)
  - $T_2$  grabs locks  $L_1$  and  $L_2$  (in some order)
  - $T_3$  grabs locks  $L_2$
  - $T_4$  grabs no locks

# Deadlock Avoidance: Example

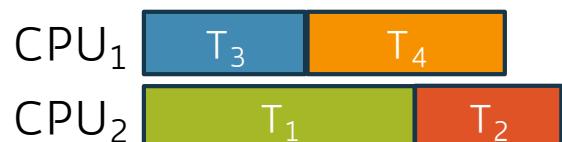
- 4 threads:  $T_1, T_2, T_3, T_4$
- 2 CPUs:  $CPU_1, CPU_2$
- knowledge:  $T_1$  grabs locks  $L_1$  and  $L_2$  (in some order)  
 $T_2$  grabs locks  $L_1$  and  $L_2$  (in some order)  
 $T_3$  grabs locks  $L_2$   
 $T_4$  grabs no locks

A smart scheduler can avoid deadlock by not running  $T_1$  and  $T_2$  in parallel

# Deadlock Avoidance: Example

- 4 threads:  $T_1, T_2, T_3, T_4$
- 2 CPUs:  $CPU_1, CPU_2$
- knowledge:  $T_1$  grabs locks  $L_1$  and  $L_2$  (in some order)  
 $T_2$  grabs locks  $L_1$  and  $L_2$  (in some order)  
 $T_3$  grabs locks  $L_2$   
 $T_4$  grabs no locks

A smart scheduler can avoid deadlock by not running  $T_1$  and  $T_2$  in parallel



$T_3$  and  $T_1$  ( $T_2$ ) can run in parallel

# Deadlock Avoidance: Safe State

- There exists **at least one** sequence of execution (safe sequence)

# Deadlock Avoidance: Safe State

- There exists **at least one** sequence of execution (safe sequence)
- Each thread can obtain all the resources it needs, complete its execution, and release the resources

# Deadlock Avoidance: Safe State

- There exists **at least one** sequence of execution (safe sequence)
- Each thread can obtain all the resources it needs, complete its execution, and release the resources
- An unsafe state does not necessarily mean deadlock (i.e., some threads may not request all the resources they declared)

# Deadlock Avoidance: Safe State

- There exists **at least one** sequence of execution (safe sequence)
- Each thread can obtain all the resources it needs, complete its execution, and release the resources
- An unsafe state does not necessarily mean deadlock (i.e., some threads may not request all the resources they declared)
- Grant a resource to a thread if the new state is safe, otherwise make it wait even if the resource is available

# Deadlock Avoidance: Safe State

- There exists **at least one** sequence of execution (safe sequence)
- Each thread can obtain all the resources it needs, complete its execution, and release the resources
- An unsafe state does not necessarily mean deadlock (i.e., some threads may not request all the resources they declared)
- Grant a resource to a thread if the new state is safe, otherwise make it wait even if the resource is available
- Given  $n$  threads, the naïve, brute-force approach would require analyzing **all the possible permutations** of them  $O(n!)$

# Banker's Algorithm

- Handles multiple instances of the same resource
- Forces threads to provide information on what resource they might need, in advance
- The resources requested must not exceed the total available in the system
- The algorithm allocates resources to a requesting thread if the allocation leaves the system in a safe state, otherwise the thread waits

# Banker's Algorithm: Idea

- Keep track of which threads can finish given the current available resources

# Banker's Algorithm: Idea

- Keep track of which threads can finish given the current available resources
- Whenever one thread that can finish is found, pretend it finishes and releases its resources

# Banker's Algorithm: Idea

- Keep track of which threads can finish given the current available resources
- Whenever one thread that can finish is found, pretend it finishes and releases its resources
- Repeat the process until:
  - all processes can finish ( $\rightarrow$  **safe state**), or
  - no remaining process can finish ( $\rightarrow$  **unsafe state**)

# Banker's Algorithm: Idea

- Keep track of which threads can finish given the current available resources
- Whenever one thread that can finish is found, pretend it finishes and releases its resources
- Repeat the process until:
  - all processes can finish ( $\rightarrow$  **safe state**), or
  - no remaining process can finish ( $\rightarrow$  **unsafe state**)
- This solution ensures that if any safe sequence exists, one will be found **without needing to test all permutations**

# Banker's Algorithm: Data Structures

- $n$  = number of threads;  $m$  = number of resource types
- $\text{available}[1..m]$ :  $m$ -dimensional vector
  - $\text{available}[j] = k$  means there are  $k$  resources of type  $j$  available
- $\text{max}[1..n, 1..m]$ :  $n \times m$  matrix
  - $\text{max}[i, j] = k$  means thread  $i$  may require at most  $k$  resources of type  $j$
- $\text{allocation}[1..n, 1..m]$ :  $n \times m$  matrix
  - $\text{allocation}[i, j] = k$  means thread  $i$  has allocated  $k$  resources of type  $j$
- $\text{need}[1..n, 1..m]$ :  $n \times m$  matrix
  - $\text{need}[i, j] = \text{max}[i, j] - \text{allocation}[i, j] = k$  means thread  $i$  may need  $k$  more resources of type  $j$  to complete its task

# Banker's Algorithm: Pseudocode

```
work = available                                // m-dimensional vectors
finish[i] = false for all i

repeat                                              // outer loop
    found = false
    for i in 1..n:                               // inner loop
        if not finish[i] and need[i] <= work:
            work = work + allocation[i]
            finish[i] = true
            found = true
    until not found

    if all finish[i] == true // safe state
        return true
    return false                                // unsafe state
```

# Banker's Algorithm: Time Complexity

```
work = available                      // m-dimensional vectors
finish[i] = false for all i

repeat                                     // outer loop
    found = false
    for i in 1..n:                         // inner loop
        if not finish[i] and need[i] <= work:
            work = work + allocation[i]
            finish[i] = true
            found = true
until not found

if all finish[i] == true // safe state
    return true
return false                                // unsafe state
```

The **inner loop** scans all n threads to look for one that can finish → O(n)

# Banker's Algorithm: Time Complexity

```
work = available                                // m-dimensional vectors
finish[i] = false for all i

repeat                                         // outer loop
    found = false
    for i in 1..n:                               // inner loop
        if not finish[i] and need[i] <= work:
            work = work + allocation[i]
            finish[i] = true
            found = true
until not found

if all finish[i] == true // safe state
    return true
return false                                     // unsafe state
```

The **outer loop** runs as long as we find at least one process that can finish.

In the worst case, finding and completing only one thread at each iteration →  $O(n)$

# Banker's Algorithm: Time Complexity

```
work = available          // m-dimensional vectors
finish[i] = false for all i

repeat
    found = false
    for i in 1..n:           // inner loop
        if not finish[i] and need[i] <= work:
            work = work + allocation[i]
            finish[i] = true
            found = true
until not found

if all finish[i] == true // safe state
    return true
return false             // unsafe state
```

Checking  $\text{need}[i] \leq \text{work}$  may require up to  $m$  resource types →  $O(m)$  per thread

# Banker's Algorithm: Time Complexity

```
work = available                      // m-dimensional vectors
finish[i] = false for all i

repeat                                // outer loop
    found = false
    for i in 1..n:                  // inner loop
        if not finish[i] and need[i] <= work:
            work = work + allocation[i]
            finish[i] = true
            found = true
    until not found

    if all finish[i] == true // safe state
        return true
    return false                   // unsafe state
```

Overall,  
 $O(n^2*m)$

# Banker's Algorithm: Example

A snapshot of the current state of the system

THREADS		RESOURCES								
		MAX			ALLOCATION			AVAILABLE		
		A	B	C	A	B	C	A	B	C
T <sub>0</sub>		0	0	1	0	0	1			
T <sub>1</sub>		1	7	5	1	0	0			
T <sub>2</sub>		2	3	5	1	3	5			
T <sub>3</sub>		0	6	5	0	6	3			
Total					2	9	9	1	5	2

# Banker's Algorithm: Example

Q1: How many resources of type A, B, and C are there overall?

		RESOURCES								
T H R E A D S	T <sub>0</sub>	MAX			ALLOCATION			AVAILABLE		
		A	B	C	A	B	C	A	B	C
T <sub>0</sub>	0	0	1	0	0	1				
T <sub>1</sub>	1	7	5	1	0	0				
T <sub>2</sub>	2	3	5	1	3	5				
T <sub>3</sub>	0	6	5	0	6	3				
Total				2	9	9	1	5	2	

# Banker's Algorithm: Example

Q1: How many resources of type A, B, and C are there overall?

THREADS		RESOURCES								
		MAX			ALLOCATION			AVAILABLE		
		A	B	C	A	B	C	A	B	C
T <sub>0</sub>		0	0	1	0	0	1			
T <sub>1</sub>		1	7	5	1	0	0			
T <sub>2</sub>		2	3	5	1	3	5			
T <sub>3</sub>		0	6	5	0	6	3			
Total					2	9	9	1	5	2

$$\begin{aligned}A &= 2 + 1 = 3 \\B &= 9 + 5 = 14 \\C &= 9 + 2 = 11\end{aligned}$$

# Banker's Algorithm: Example

Q2: What is the content of the NEED matrix?

		RESOURCES											
THREADS	T <sub>0</sub>	MAX			ALLOCATION		AVAILABLE			NEED			
		A	B	C	A	B	C	A	B	C	A	B	C
	T <sub>1</sub>	1	7	5	1	0	0	1	2	6	5	2	1
	T <sub>2</sub>	2	3	5	1	3	5	1	5	2	1	3	5
	T <sub>3</sub>	0	6	5	0	6	3	2	9	9	0	3	2
	Total				2	9	9	1	5	2			

# Banker's Algorithm: Example

Q2: What is the content of the NEED matrix?

$$\text{NEED}[i, j] = \text{MAX}[i, j] - \text{ALLOCATION}[i, j]$$

		RESOURCES									NEED		
		MAX			ALLOCATION			AVAILABLE					
		A	B	C	A	B	C	A	B	C	A	B	C
THREADS	T <sub>0</sub>	0	0	1	0	0	1						
	T <sub>1</sub>	1	7	5	1	0	0						
	T <sub>2</sub>	2	3	5	1	3	5						
	T <sub>3</sub>	0	6	5	0	6	3						
	Total				2	9	9	1	5	2			

# Banker's Algorithm: Example

Q2: What is the content of the NEED matrix?

$$\text{NEED}[i, j] = \text{MAX}[i, j] - \text{ALLOCATION}[i, j]$$

		RESOURCES						NEED			
T H R E A D S	T <sub>0</sub>	MAX			ALLOCATION			AVAILABLE			A 0-0 = 0
		A	B	C	A	B	C	A	B	C	
		0	0	1	0	0	1	2	9	9	
		1	7	5	1	0	0	1	5	2	
		2	3	5	1	3	5	1	5	2	
		0	6	5	0	6	3	1	5	2	
Total											

# Banker's Algorithm: Example

Q2: What is the content of the NEED matrix?

$$\text{NEED}[i, j] = \text{MAX}[i, j] - \text{ALLOCATION}[i, j]$$

		RESOURCES									NEED		
T H R E A D S	T <sub>0</sub>	MAX			ALLOCATION			AVAILABLE			A	B	C
		0	0	1	0	0	1	A	B	C	0	0-0 =	0
		1	7	5	1	0	0	A	B	C	1	6	5
		2	3	5	1	3	5	A	B	C	1	5	2
	T <sub>3</sub>	0	6	5	0	6	3	A	B	C	1	5	2
Total					2	9	9	1	5	2			

# Banker's Algorithm: Example

Q2: What is the content of the NEED matrix?

$$\text{NEED}[i, j] = \text{MAX}[i, j] - \text{ALLOCATION}[i, j]$$

		RESOURCES									NEED		
THREADS	T <sub>0</sub>	MAX			ALLOCATION			AVAILABLE			A	B	C
		A	B	C	A	B	C	A	B	C	A	B	C
T <sub>0</sub>	T <sub>0</sub>	0	0	1	0	0	1	1	0	0	0	0	1-1 = 0
T <sub>1</sub>	T <sub>1</sub>	1	7	5	1	0	0	1	5	2	0	2	3
T <sub>2</sub>	T <sub>2</sub>	2	3	5	1	3	5	2	3	2	0	1	4
T <sub>3</sub>	T <sub>3</sub>	0	6	5	0	6	3	1	5	2	0	1	2
	Total				2	9	9	1	5	2			

# Banker's Algorithm: Example

Q2: What is the content of the NEED matrix?

$$\text{NEED}[i, j] = \text{MAX}[i, j] - \text{ALLOCATION}[i, j]$$

		RESOURCES									NEED		
		MAX			ALLOCATION			AVAILABLE					
		A	B	C	A	B	C	A	B	C	A	B	C
T H R E A D S	T <sub>0</sub>	0	0	1	0	0	1				0	0	0
	T <sub>1</sub>	1	7	5	1	0	0				0	7	5
	T <sub>2</sub>	2	3	5	1	3	5				1	0	0
	T <sub>3</sub>	0	6	5	0	6	3				0	0	2
	Total				2	9	9	1	5	2			

# Banker's Algorithm: Example

Q3: Is the system in a safe state? Why?

THREADS		RESOURCES									NEED					
		MAX			ALLOCATION			AVAILABLE								
		A	B	C	A	B	C	A	B	C						
T <sub>0</sub>	0	0	1	0	0	1	2	9	9	1	0	0	0			
T <sub>1</sub>	1	7	5	1	0	0	0	7	5	0	0	7	5			
T <sub>2</sub>	2	3	5	1	3	5	1	0	0	1	0	0	0			
T <sub>3</sub>	0	6	5	0	6	3	0	2	2	0	0	0	2			
Total																

# Banker's Algorithm: Example

Let's start with  $T_0$

THREADS		RESOURCES						NEED					
		MAX			ALLOCATION								
		A	B	C	A	B	C	A	B	C	A	B	C
	$T_0$	0	0	1	0	0	1				0	0	0
	$T_1$	1	7	5	1	0	0				0	7	5
	$T_2$	2	3	5	1	3	5				1	0	0
	$T_3$	0	6	5	0	6	3				0	0	2
	Total				2	9	9	1	5	2			

# Banker's Algorithm: Example

Eventually,  $T_0$  finishes and releases all its resources

THREADS		RESOURCES						NEED					
		MAX			ALLOCATION								
		A	B	C	A	B	C	A	B	C	A	B	C
T <sub>0</sub>		0	0	1	0	0	1	0	0	0	0	0	0
T <sub>1</sub>		1	7	5	1	0	0	0	7	5	0	7	5
T <sub>2</sub>		2	3	5	1	3	5	1	0	0	1	0	0
T <sub>3</sub>		0	6	5	0	6	3	0	0	2	0	0	2
Total		2	9	9	2	9	9	1	5	2	0	0	0

# Banker's Algorithm: Example

$T_1$  can't execute as it still might NEED  $(0, 7, 5)$  and AVAILABLE =  $(1, 5, 3)$

THREADS		RESOURCES						NEED		
		MAX			ALLOCATION					
		A	B	C	A	B	C	A	B	C
	$T_0$	0	0	1	-	-	-	-	-	-
	$T_1$	1	7	5	1	0	0	0	7	5
	$T_2$	2	3	5	1	3	5	1	0	0
	$T_3$	0	6	5	0	6	3	0	0	2
	Total				2	9	8	1	5	3

# Banker's Algorithm: Example

$T_2$  can execute as it still might NEED  $(1, 0, 0)$  and AVAILABLE =  $(1, 5, 3)$

THREADS		RESOURCES									NEED		
		MAX			ALLOCATION			AVAILABLE					
		A	B	C	A	B	C	A	B	C	A	B	C
	$T_0$	0	0	1	-	-	-				-	-	-
	$T_1$	1	7	5	1	0	0				0	7	5
	$T_2$	2	3	5	1	3	5				1	0	0
	$T_3$	0	6	5	0	6	3				0	0	2
	Total				2	9	8	1	5	3			

# Banker's Algorithm: Example

$T_2$  can execute as it still might NEED  $(1, 0, 0)$  and AVAILABLE =  $(1, 5, 3)$

THREADS		RESOURCES									NEED		
		MAX			ALLOCATION			AVAILABLE					
		A	B	C	A	B	C	A	B	C	A	B	C
	$T_0$	0	0	1	-	-	-				-	-	-
	$T_1$	1	7	5	1	0	0				0	7	5
	$T_2$	2	3	5	2	3	5				0	0	0
	$T_3$	0	6	5	0	6	3				0	0	2
	Total				3	9	8	0	5	3			

# Banker's Algorithm: Example

$T_2$  eventually finishes and releases all its resources

THREADS		RESOURCES									NEED		
		MAX			ALLOCATION			AVAILABLE					
		A	B	C	A	B	C	A	B	C	A	B	C
$T_0$		0	0	1	-	-	-	-	-	-	-	-	-
$T_1$		1	7	5	1	0	0	-	-	-	0	7	5
$T_2$		2	3	5	-	-	-	-	-	-	-	-	-
$T_3$		0	6	5	0	6	3	-	-	-	0	0	2
Total		1	6	3	2	8	8						

# Banker's Algorithm: Example

$T_3$  can execute as it still might NEED  $(0, 0, 2)$  and AVAILABLE =  $(2, 8, 8)$

THREADS		RESOURCES									NEED		
		MAX			ALLOCATION			AVAILABLE					
		A	B	C	A	B	C	A	B	C	A	B	C
	$T_0$	0	0	1	-	-	-				-	-	-
	$T_1$	1	7	5	1	0	0				0	7	5
	$T_2$	2	3	5	-	-	-				-	-	-
	$T_3$	0	6	5	0	6	3				0	0	2
	Total				1	6	3	2	8	8			

# Banker's Algorithm: Example

$T_3$  can execute as it still might NEED  $(0, 0, 2)$  and AVAILABLE =  $(2, 3, 6)$

THREADS		RESOURCES						NEED					
		MAX			ALLOCATION								
		A	B	C	A	B	C	A	B	C	A	B	C
	$T_0$	0	0	1	-	-	-	-	-	-	-	-	-
	$T_1$	1	7	5	1	0	0	0	7	5	0	7	5
	$T_2$	2	3	5	-	-	-	-	-	-	-	-	-
	$T_3$	0	6	5	0	6	5	2	8	6	0	0	0
	Total				1	6	5	2	8	6			

# Banker's Algorithm: Example

$T_3$  eventually finishes and releases all its resources

THREADS		RESOURCES									NEED		
		MAX			ALLOCATION			AVAILABLE					
		A	B	C	A	B	C	A	B	C	A	B	C
T <sub>0</sub>		0	0	1	-	-	-	-	-	-	-	-	-
T <sub>1</sub>		1	7	5	1	0	0	-	-	-	0	7	5
T <sub>2</sub>		2	3	5	-	-	-	-	-	-	-	-	-
T <sub>3</sub>		0	6	5	-	-	-	-	-	-	-	-	-
Total		1	0	0	2	14	11						

# Banker's Algorithm: Example

$T_1$  can now execute since NEED (0, 7, 5) and AVAILABLE = (2, 14, 11)

THREADS		RESOURCES						NEED					
		MAX			ALLOCATION								
		A	B	C	A	B	C	A	B	C	A	B	C
	$T_0$	0	0	1	-	-	-	-	-	-	-	-	-
	$T_1$	1	7	5	1	7	5	0	0	0			
	$T_2$	2	3	5	-	-	-	-	-	-	-	-	-
	$T_3$	0	6	5	-	-	-	-	-	-	-	-	-
	Total				1	7	5	2	7	6			

# Banker's Algorithm: Example

We have found a sequence of execution  $T_0, T_2, T_3, T_1$   
which leads to safe state!

THREADS		RESOURCES									NEED		
		MAX			ALLOCATION			AVAILABLE					
		A	B	C	A	B	C	A	B	C	A	B	C
T <sub>0</sub>		0	0	1	-	-	-	-	-	-	-	-	-
T <sub>1</sub>		1	7	5	-	-	-	-	-	-	-	-	-
T <sub>2</sub>		2	3	5	-	-	-	-	-	-	-	-	-
T <sub>3</sub>		0	6	5	-	-	-	-	-	-	-	-	-
Total					-	-	-	3	14	11			

# Banker's Algorithm: Example

Q4: If  $T_1$  issues a REQUEST (0, 5, 2), can this be granted immediately?

THREADS		RESOURCES						NEED					
		MAX			ALLOCATION								
		A	B	C	A	B	C	A	B	C	A	B	C
T <sub>0</sub>		0	0	1	0	0	1				0	0	0
T <sub>1</sub>		1	7	5	1	0	0				0	7	5
T <sub>2</sub>		2	3	5	1	3	5				1	0	0
T <sub>3</sub>		0	6	5	0	6	3				0	0	2
Total					2	9	9	1	5	2			

# Banker's Algorithm: Example

We have to ask ourselves:

1. if the request can be satisfied;
2. if it will lead to a safe state

		RESOURCES											
T H R E A D S	T <sub>0</sub>	MAX			ALLOCATION			AVAILABLE			NEED		
		A	B	C	A	B	C	A	B	C	A	B	C
T <sub>0</sub>	T <sub>0</sub>	0	0	1	0	0	1	5	4	3	0	0	0
T <sub>1</sub>	T <sub>1</sub>	1	7	5	1	0	0	4	3	3	0	7	5
T <sub>2</sub>	T <sub>2</sub>	2	3	5	1	3	5	3	2	2	1	0	0
T <sub>3</sub>	T <sub>3</sub>	0	6	5	0	6	3	2	1	2	0	0	2
Total					2	9	9	1	5	2			

# Banker's Algorithm: Example

To answer 1. check if: a. REQUEST <= NEED and b. REQUEST <= AVAILABLE

THREADS		RESOURCES						NEED					
		MAX			ALLOCATION								
		A	B	C	A	B	C	A	B	C	A	B	C
T <sub>0</sub>		0	0	1	0	0	1				0	0	0
T <sub>1</sub>		1	7	5	1	0	0				0	7	5
T <sub>2</sub>		2	3	5	1	3	5				1	0	0
T <sub>3</sub>		0	6	5	0	6	3				0	0	2
	Total				2	9	9	1	5	2			

# Banker's Algorithm: Example

1.a. REQUEST <= NEED?

		RESOURCES									NEED		
		MAX			ALLOCATION			AVAILABLE					
		A	B	C	A	B	C	A	B	C	A	B	C
T H R E A D S	T <sub>0</sub>	0	0	1	0	0	1				0	0	0
	T <sub>1</sub>	1	7	5	1	0	0				0	7	5
	T <sub>2</sub>	2	3	5	1	3	5				1	0	0
	T <sub>3</sub>	0	6	5	0	6	3				0	0	2
Total					2	9	9	1	5	2			

# Banker's Algorithm: Example

1.a. REQUEST  $\leq$  NEED? YES!  $(0, 5, 2) \leq (0, 7, 5)$

		RESOURCES									NEED					
		MAX			ALLOCATION			AVAILABLE								
		A	B	C	A	B	C	A	B	C						
T H R E A D S	T <sub>0</sub>	0	0	1	0	0	1				0	0	0			
	T <sub>1</sub>	1	7	5	1	0	0				0	7	5			
	T <sub>2</sub>	2	3	5	1	3	5				1	0	0			
	T <sub>3</sub>	0	6	5	0	6	3				0	0	2			
Total					2	9	9	1	5	2						

# Banker's Algorithm: Example

1.b. REQUEST <= AVAILABLE?

		RESOURCES									NEED					
		MAX			ALLOCATION			AVAILABLE								
		A	B	C	A	B	C	A	B	C						
T H R E A D S	T <sub>0</sub>	0	0	1	0	0	1				0	0	0			
	T <sub>1</sub>	1	7	5	1	0	0				0	7	5			
	T <sub>2</sub>	2	3	5	1	3	5				1	0	0			
	T <sub>3</sub>	0	6	5	0	6	3				0	0	2			
Total					2	9	9	1	5	2						

# Banker's Algorithm: Example

1.b. REQUEST  $\leq$  AVAILABLE? YES!  $(0, 5, 2) \leq (1, 5, 2)$

THREADS		RESOURCES									NEED					
		MAX			ALLOCATION			AVAILABLE								
		A	B	C	A	B	C	A	B	C						
T <sub>0</sub>	0	0	1	0	0	1	1	5	2	0	0	0	0			
T <sub>1</sub>	1	7	5	1	0	0	0	0	7	5	0	7	5			
T <sub>2</sub>	2	3	5	1	3	5	5	5	0	0	1	0	0			
T <sub>3</sub>	0	6	5	0	6	3	3	2	9	9	0	0	2			
Total					2	9	9	1	5	2						

# Banker's Algorithm: Example

To answer 2. we simulate the request is granted  
and see if we are still in a safe state

THREADS		RESOURCES									NEED		
		MAX			ALLOCATION			AVAILABLE					
		A	B	C	A	B	C	A	B	C	A	B	C
T <sub>0</sub>		0	0	1	0	0	1				0	0	0
T <sub>1</sub>		1	7	5	1	0	0				0	7	5
T <sub>2</sub>		2	3	5	1	3	5				1	0	0
T <sub>3</sub>		0	6	5	0	6	3				0	0	2
Total					2	9	9	1	5	2			

# Banker's Algorithm: Example

To answer 2. we simulate the request is granted  
and see if we are still in a safe state

THREADS	T <sub>0</sub>	MAX			ALLOCATION			AVAILABLE			NEED		
		A	B	C	A	B	C	A	B	C	A	B	C
		0	0	1	0	0	1	0	0	0	0	0	0
T <sub>1</sub>	1	7	5	1	5	2	0	2	3	0	2	3	0
T <sub>2</sub>	2	3	5	1	3	5	1	0	0	1	0	0	0
T <sub>3</sub>	0	6	5	0	6	3	0	0	2	0	0	0	2
Total				2	14	11	1	0	0				

# Banker's Algorithm: Example

Let's start with  $T_0$

THREADS		RESOURCES						NEED					
		MAX			ALLOCATION								
		A	B	C	A	B	C	A	B	C	A	B	C
	$T_0$	0	0	1	0	0	1				0	0	0
	$T_1$	1	7	5	1	5	2				0	2	3
	$T_2$	2	3	5	1	3	5				1	0	0
	$T_3$	0	6	5	0	6	3				0	0	2
	Total				2	14	11	1	0	0			

# Banker's Algorithm: Example

Eventually,  $T_0$  finishes and releases all its resources

THREADS		RESOURCES						NEED					
		MAX			ALLOCATION								
		A	B	C	A	B	C	A	B	C	A	B	C
T <sub>0</sub>		0	0	1	-	-	-	-	-	-	-	-	-
T <sub>1</sub>		1	7	5	1	5	2	0	2	3			
T <sub>2</sub>		2	3	5	1	3	5	1	0	0			
T <sub>3</sub>		0	6	5	0	6	3	0	0	2			
Total					2	14	10	1	0	1			

# Banker's Algorithm: Example

$T_1$  can't execute as it still might NEED  $(0, 2, 3)$  and AVAILABLE =  $(1, 0, 1)$

THREADS		RESOURCES						NEED		
		MAX			ALLOCATION					
		A	B	C	A	B	C	A	B	C
	$T_0$	0	0	1	-	-	-	-	-	-
	$T_1$	1	7	5	1	5	2	0	2	3
	$T_2$	2	3	5	1	3	5	1	0	0
	$T_3$	0	6	5	0	6	3	0	0	2
	Total				2	14	10	1	0	1

# Banker's Algorithm: Example

$T_2$  can execute as it still might NEED  $(1, 0, 0)$  and AVAILABLE =  $(1, 0, 1)$

THREADS		RESOURCES						NEED					
		MAX			ALLOCATION								
		A	B	C	A	B	C	A	B	C	A	B	C
	$T_0$	0	0	1	-	-	-	-	-	-	-	-	-
	$T_1$	1	7	5	1	5	2	0	2	3			
	$T_2$	2	3	5	1	3	5	1	0	0	1	0	0
	$T_3$	0	6	5	0	6	3	0	0	2			
	Total				2	14	10	1	0	1			

# Banker's Algorithm: Example

$T_2$  can execute as it still might NEED  $(1, 0, 0)$  and AVAILABLE =  $(1, 0, 1)$

THREADS		RESOURCES									NEED		
		MAX			ALLOCATION			AVAILABLE					
		A	B	C	A	B	C	A	B	C	A	B	C
T <sub>0</sub>		0	0	1	-	-	-	-	-	-	-	-	-
T <sub>1</sub>		1	7	5	1	5	2	0	2	3	0	2	3
T <sub>2</sub>		2	3	5	2	3	5	0	0	0	0	0	0
T <sub>3</sub>		0	6	5	0	6	3	0	0	2	0	0	2
Total					3	14	10	0	0	1			

# Banker's Algorithm: Example

$T_2$  eventually finishes and releases all its resources

THREADS		RESOURCES									NEED		
		MAX			ALLOCATION			AVAILABLE					
		A	B	C	A	B	C	A	B	C	A	B	C
	$T_0$	0	0	1	-	-	-	-	-	-	-	-	-
	$T_1$	1	7	5	1	5	2	-	-	-	0	2	3
	$T_2$	2	3	5	-	-	-	-	-	-	-	-	-
	$T_3$	0	6	5	0	6	3	-	-	-	0	0	2
	Total	1	11	5	2	3	6						

# Banker's Algorithm: Example

$T_3$  can execute as it still might NEED  $(0, 0, 2)$  and AVAILABLE =  $(2, 3, 6)$

THREADS		RESOURCES						NEED					
		MAX			ALLOCATION								
		A	B	C	A	B	C	A	B	C	A	B	C
	$T_0$	0	0	1	-	-	-	-	-	-	-	-	-
	$T_1$	1	7	5	1	5	2	0	2	3			
	$T_2$	2	3	5	-	-	-	-	-	-			
	$T_3$	0	6	5	0	6	3	0	0	2			
	Total				1	11	5	2	3	6			

# Banker's Algorithm: Example

$T_3$  can execute as it still might NEED  $(0, 0, 2)$  and AVAILABLE =  $(2, 3, 6)$

THREADS		RESOURCES						NEED					
		MAX			ALLOCATION								
		A	B	C	A	B	C	A	B	C	A	B	C
T <sub>0</sub>		0	0	1	-	-	-	-	-	-	-	-	-
T <sub>1</sub>		1	7	5	1	5	2	0	2	3			
T <sub>2</sub>		2	3	5	-	-	-	-	-	-			
T <sub>3</sub>		0	6	5	0	6	5	0	0	0	0		
Total					1	11	7	2	3	4			

# Banker's Algorithm: Example

$T_3$  eventually finishes and releases all its resources

THREADS		RESOURCES									NEED		
		MAX			ALLOCATION			AVAILABLE					
		A	B	C	A	B	C	A	B	C	A	B	C
T <sub>0</sub>		0	0	1	-	-	-	-	-	-	-	-	-
T <sub>1</sub>		1	7	5	1	5	2	-	-	-	0	2	3
T <sub>2</sub>		2	3	5	-	-	-	-	-	-	-	-	-
T <sub>3</sub>		0	6	5	-	-	-	-	-	-	-	-	-
Total					1	5	2	2	9	9			

# Banker's Algorithm: Example

$T_1$  can now execute since NEED ( $0, 2, 3$ ) and AVAILABLE =  $(2, 9, 9)$

THREADS		RESOURCES									NEED		
		MAX			ALLOCATION			AVAILABLE					
		A	B	C	A	B	C	A	B	C	A	B	C
	$T_0$	0	0	1	-	-	-	-	-	-	-	-	-
	$T_1$	1	7	5	1	7	5	0	0	0			
	$T_2$	2	3	5	-	-	-	-	-	-			
	$T_3$	0	6	5	-	-	-	-	-	-			
	Total				1	7	5	2	7	6			

# Banker's Algorithm: Example

We have found a sequence of execution  $T_0, T_2, T_3, T_1$   
which leads to safe state!

THREADS		RESOURCES									NEED		
		MAX			ALLOCATION			AVAILABLE					
		A	B	C	A	B	C	A	B	C	A	B	C
T <sub>0</sub>		0	0	1	-	-	-	-	-	-	-	-	-
T <sub>1</sub>		1	7	5	-	-	-	-	-	-	-	-	-
T <sub>2</sub>		2	3	5	-	-	-	-	-	-	-	-	-
T <sub>3</sub>		0	6	5	-	-	-	-	-	-	-	-	-
Total		-	-	-	-	-	-	3	14	11			

# Our Journey

- What is deadlock?
- Conditions for deadlock to happen
- Deadlock prevention
- Deadlock avoidance
- **Deadlock detection and recovery**

# Deadlock Detection: Resource Allocation Graph

- We define a directed graph  $G=(V, E)$  where:

# Deadlock Detection: Resource Allocation Graph

- We define a **directed graph**  $G=(V, E)$  where:
  - $V$  is the set of vertices representing both **resources**  $\{r_1, \dots, r_m\}$  and **threads**  $\{t_1, \dots, t_n\}$

# Deadlock Detection: Resource Allocation Graph

- We define a **directed graph**  $G=(V, E)$  where:
  - $V$  is the set of vertices representing both **resources**  $\{r_1, \dots, r_m\}$  and **threads**  $\{t_1, \dots, t_n\}$
  - $E$  is the set of edges between resources and threads

# Deadlock Detection: Resource Allocation Graph

- We define a **directed graph**  $G=(V, E)$  where:
  - $V$  is the set of vertices representing both **resources**  $\{r_1, \dots, r_m\}$  and **threads**  $\{t_1, \dots, t_n\}$
  - $E$  is the set of edges between resources and threads
- Edges can be of **2 types**:

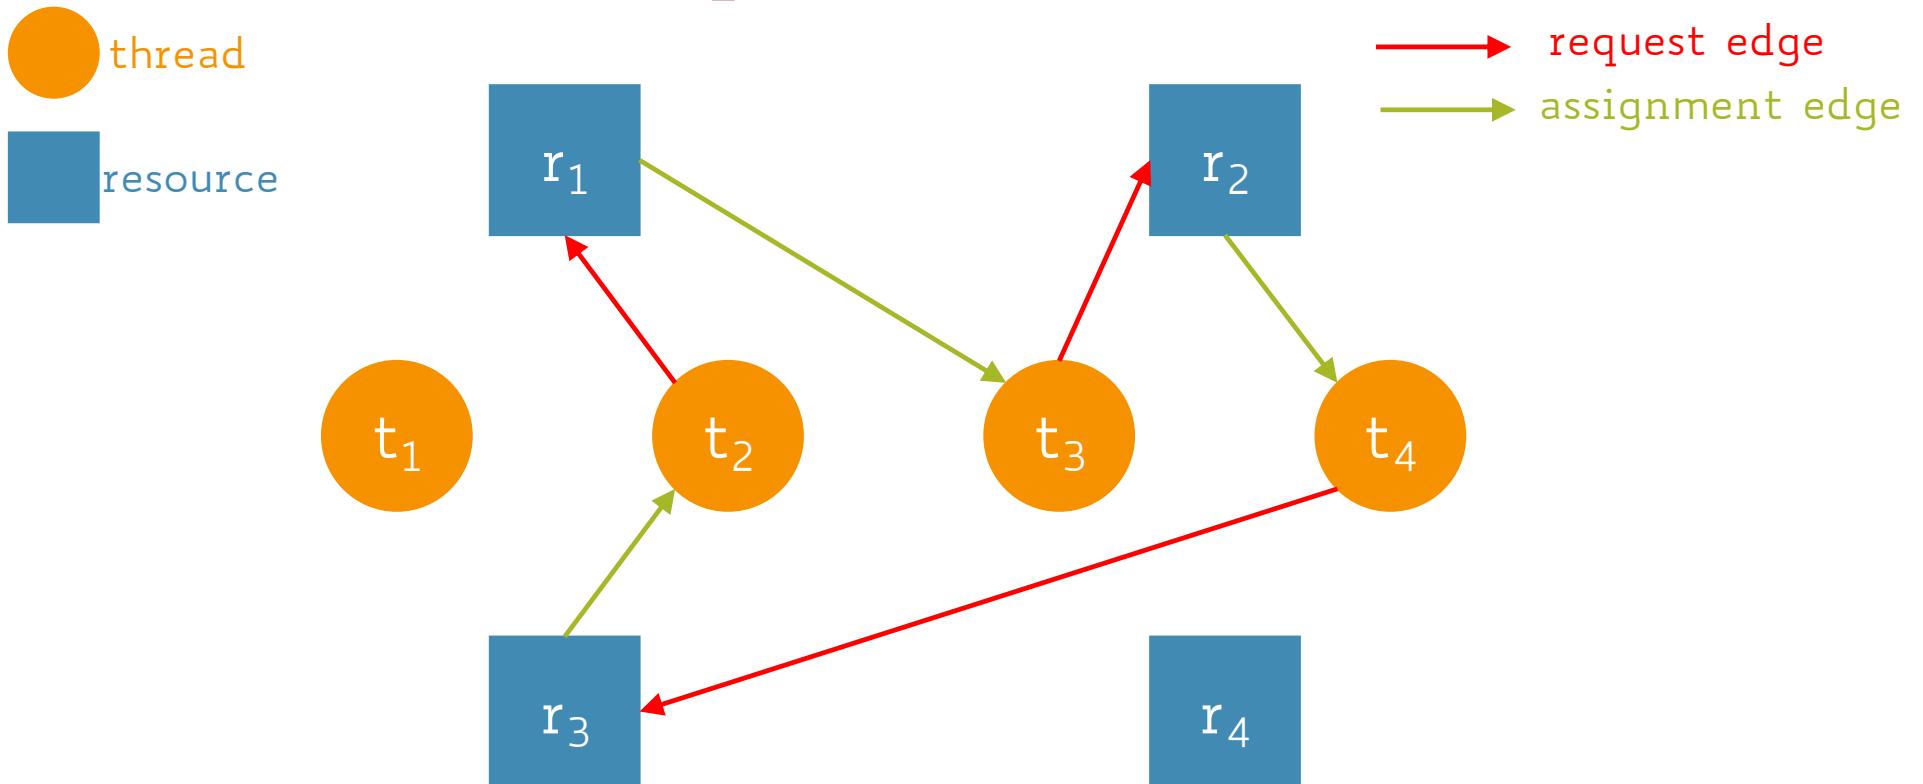
# Deadlock Detection: Resource Allocation Graph

- We define a **directed graph**  $G=(V, E)$  where:
  - $V$  is the set of vertices representing both **resources**  $\{r_1, \dots, r_m\}$  and **threads**  $\{t_1, \dots, t_n\}$
  - $E$  is the set of edges between resources and threads
- Edges can be of **2 types**:
  - **Request Edge** → a directed edge  $(t_i, r_j)$  indicates that  $t_i$  has requested  $r_j$ , but not yet acquired

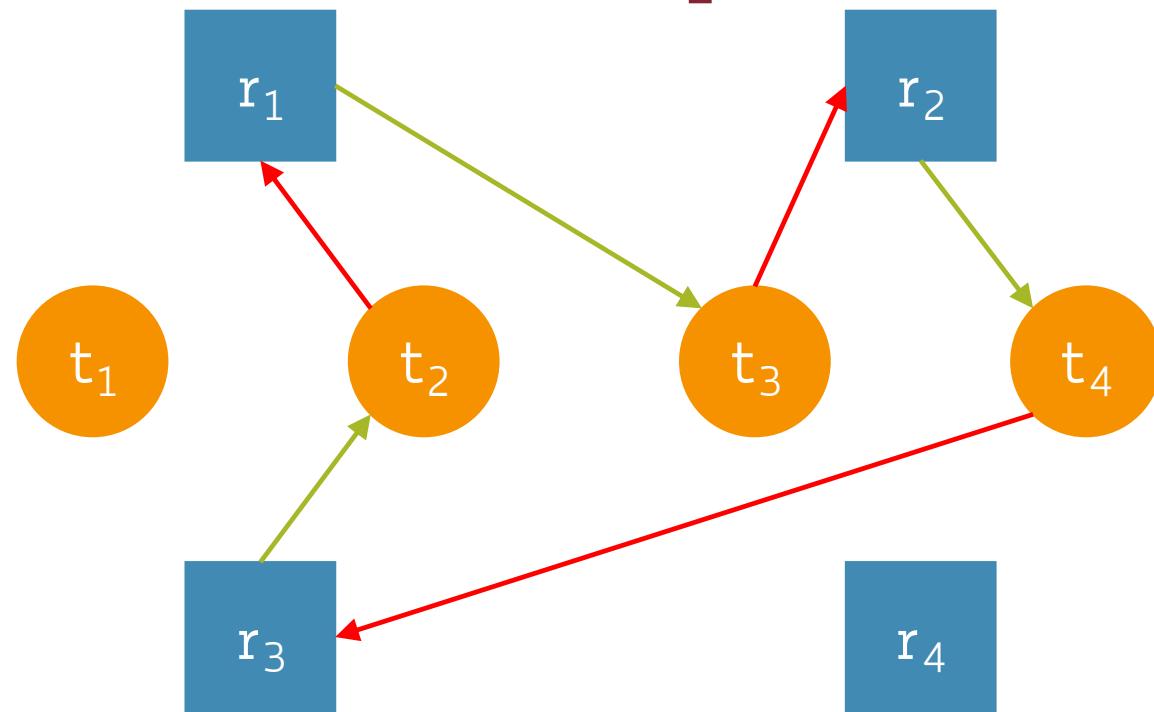
# Deadlock Detection: Resource Allocation Graph

- We define a **directed graph**  $G=(V, E)$  where:
  - $V$  is the set of vertices representing both **resources**  $\{r_1, \dots, r_m\}$  and **threads**  $\{t_1, \dots, t_n\}$
  - $E$  is the set of edges between resources and threads
- Edges can be of **2 types**:
  - **Request Edge** → a directed edge  $(t_i, r_j)$  indicates that  $t_i$  has requested  $r_j$ , but not yet acquired
  - **Assignment Edge** → a directed edge  $(r_j, t_i)$  indicates that the OS has allocated  $r_j$  to  $t_i$

# Deadlock Detection: Resource Allocation Graph

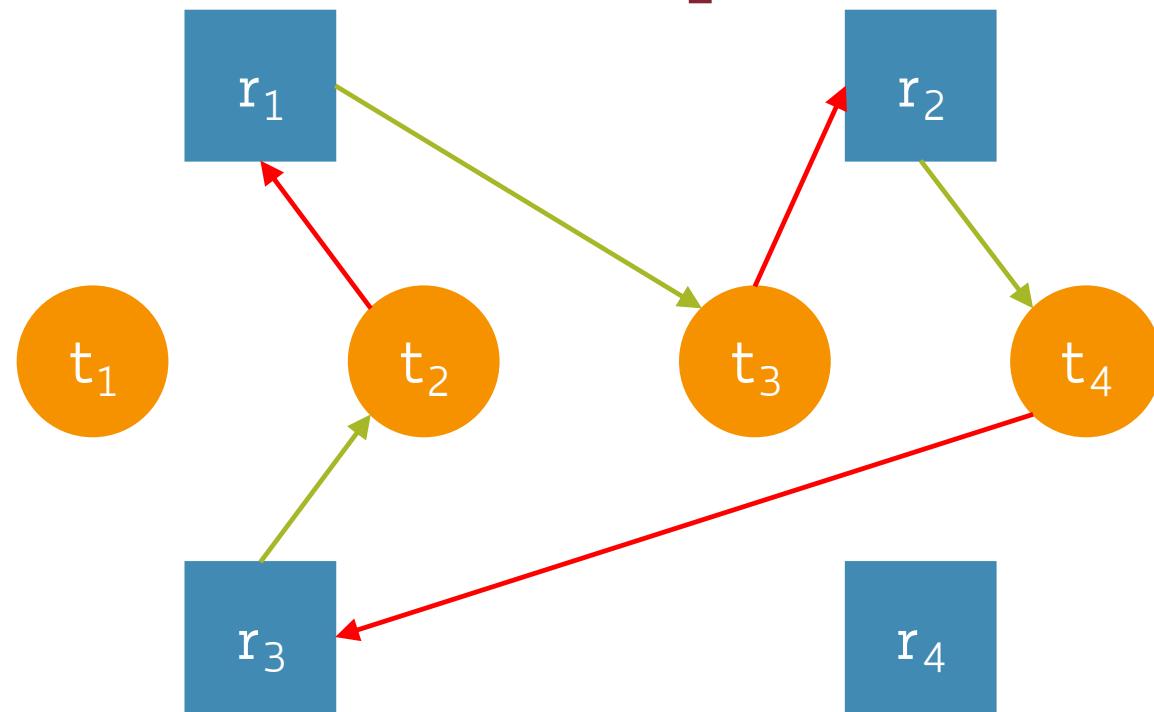


# Deadlock Detection: Resource Allocation Graph



If the graph has no cycles, no deadlock will ever exist

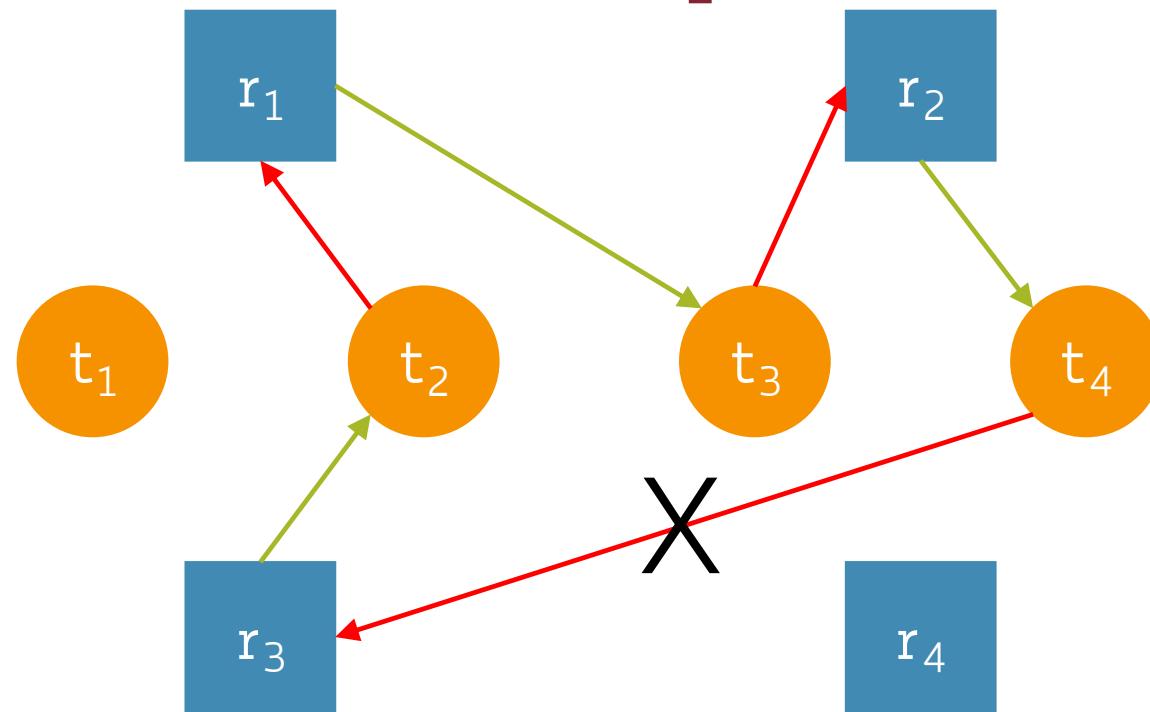
# Deadlock Detection: Resource Allocation Graph



If the graph has no cycles, no deadlock will ever exist

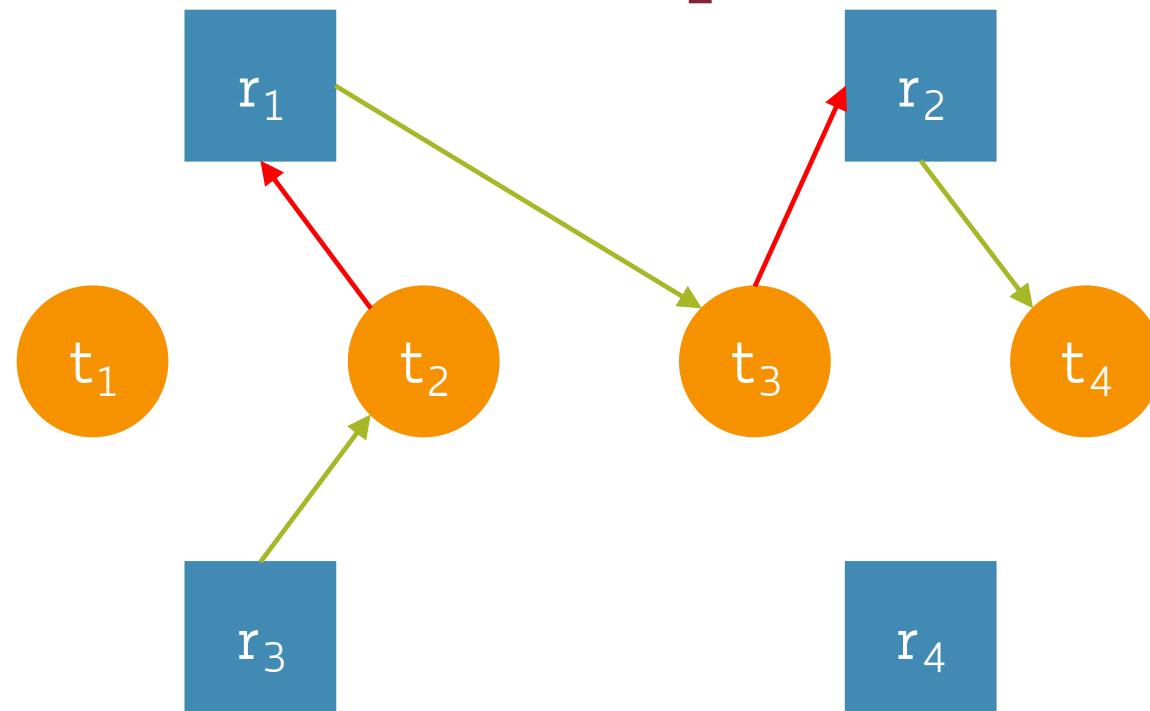
Why?

# Deadlock Detection: Resource Allocation Graph



Suppose we remove the edge  $(t_4, r_3)$  so as to remove the cycle

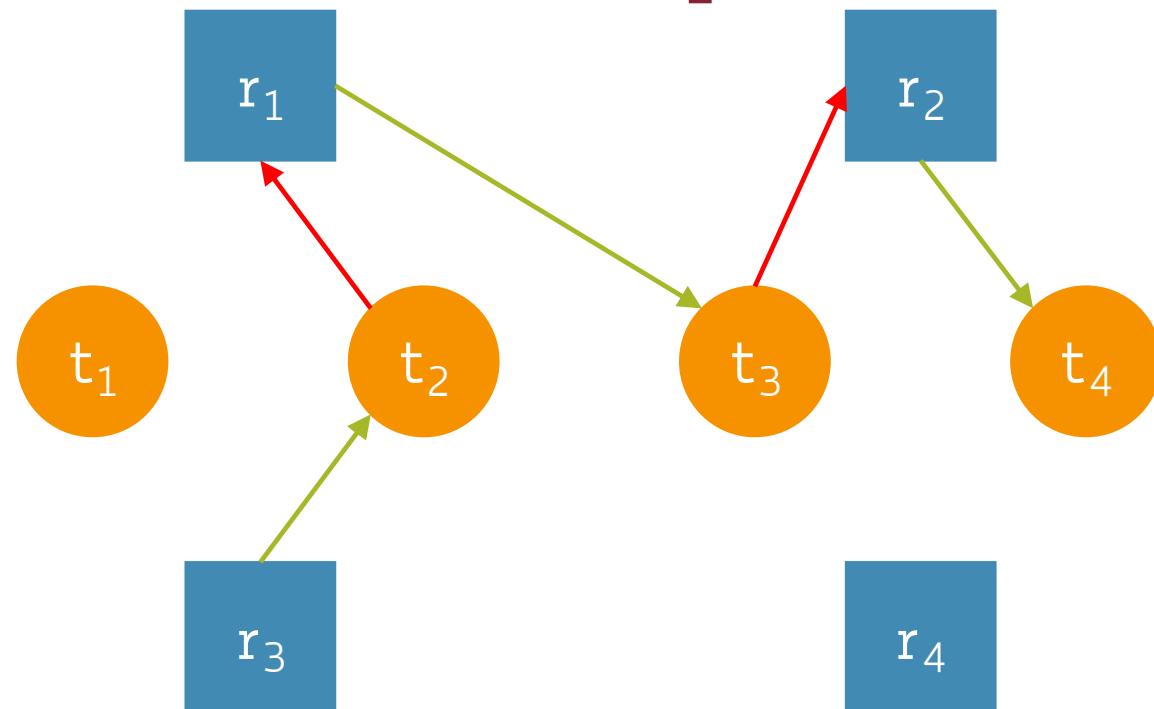
# Deadlock Detection: Resource Allocation Graph



Suppose we remove the edge  $(t_4, r_3)$  so as to remove the cycle

No deadlock can occur as  $t_4$  is not waiting on anything...

# Deadlock Detection: Resource Allocation Graph

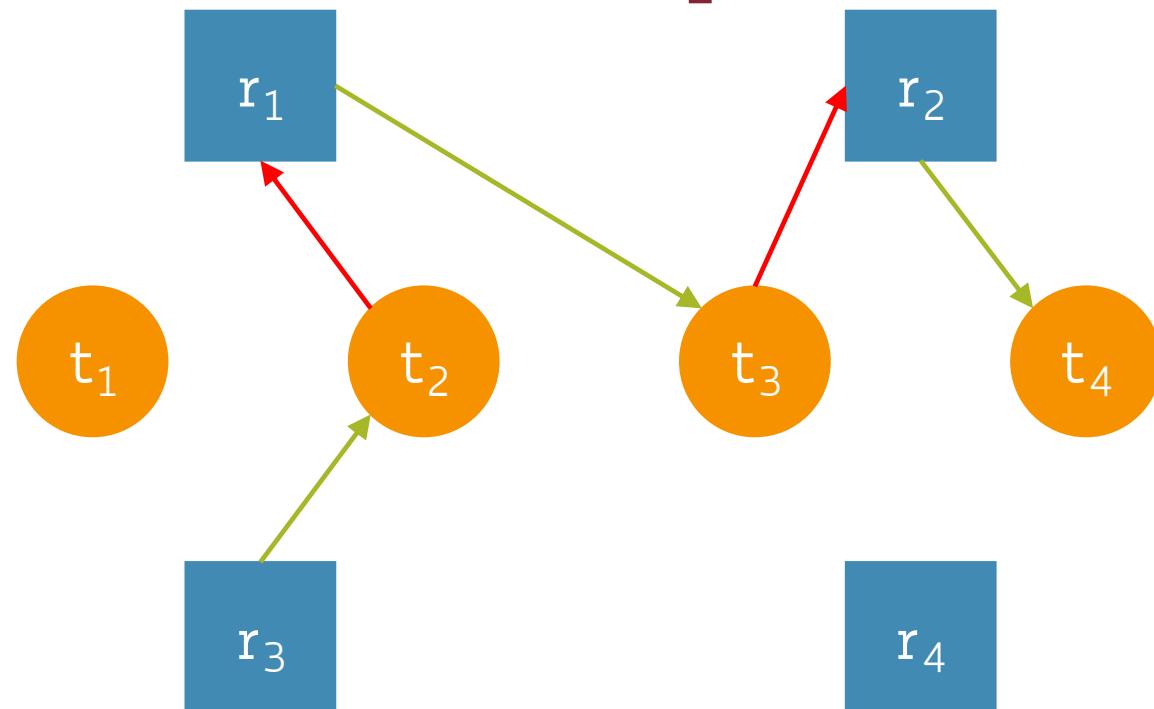


Suppose we remove the edge  $(t_4, r_3)$  so as to remove the cycle

No deadlock can occur as  $t_4$  is not waiting on anything...

Therefore,  $t_4$  can run and eventually will release  $r_2$ , which wakes up  $t_3$

# Deadlock Detection: Resource Allocation Graph




Suppose we remove the edge  $(t_4, r_3)$  so as to remove the cycle

No deadlock can occur as  $t_4$  is not waiting on anything...

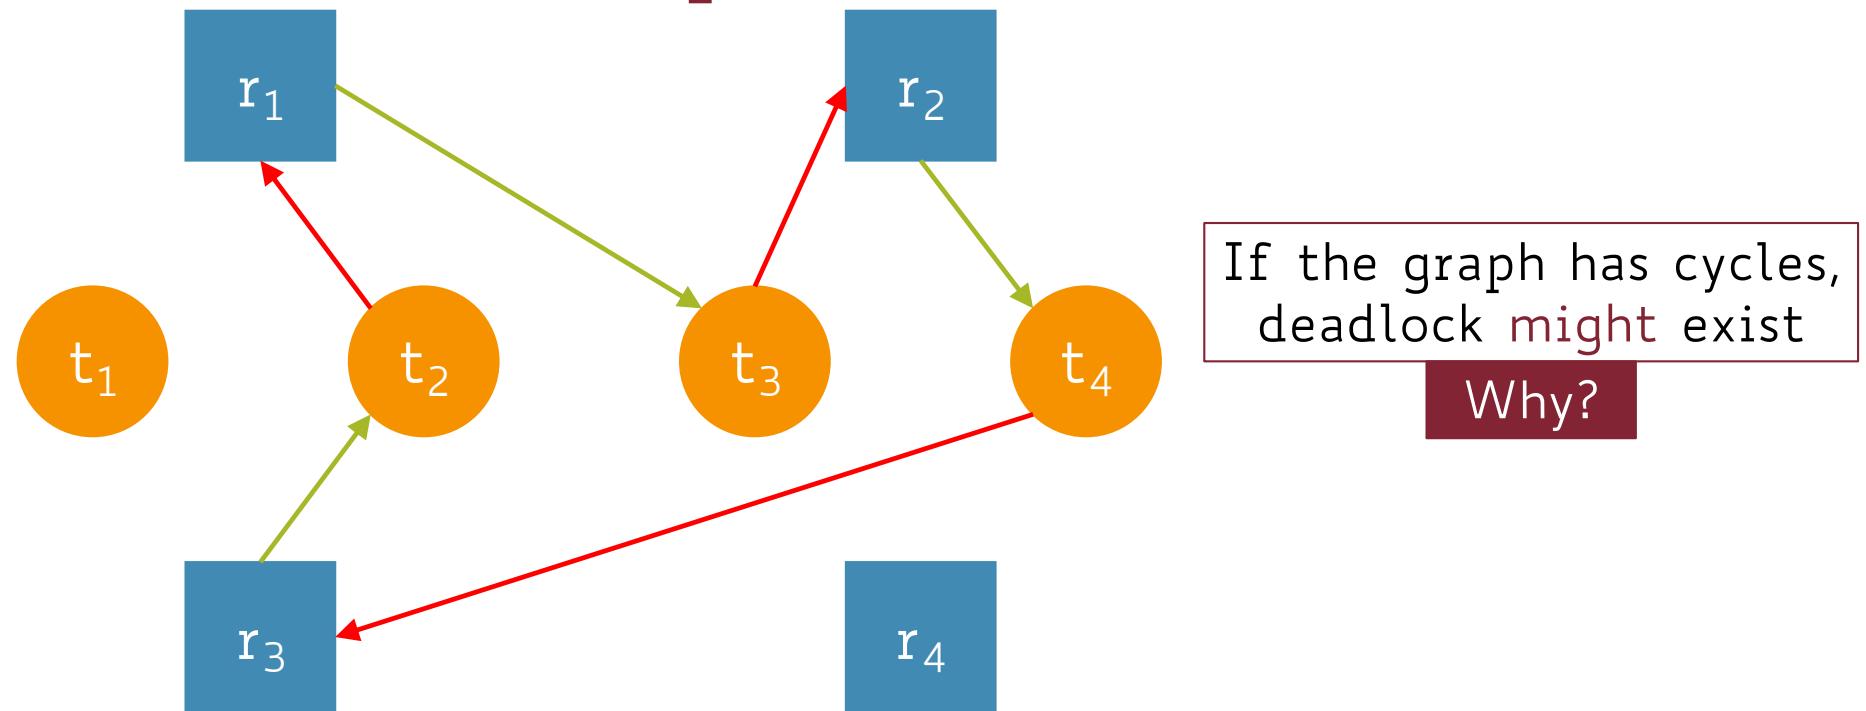
Therefore,  $t_4$  can run and eventually will release  $r_2$ , which wakes up  $t_3$

And so on and so forth...


# Deadlock Detection: Resource Allocation Graph



# Deadlock Detection: Resource Allocation Graph




# Deadlock Detection: Resource Allocation Graph



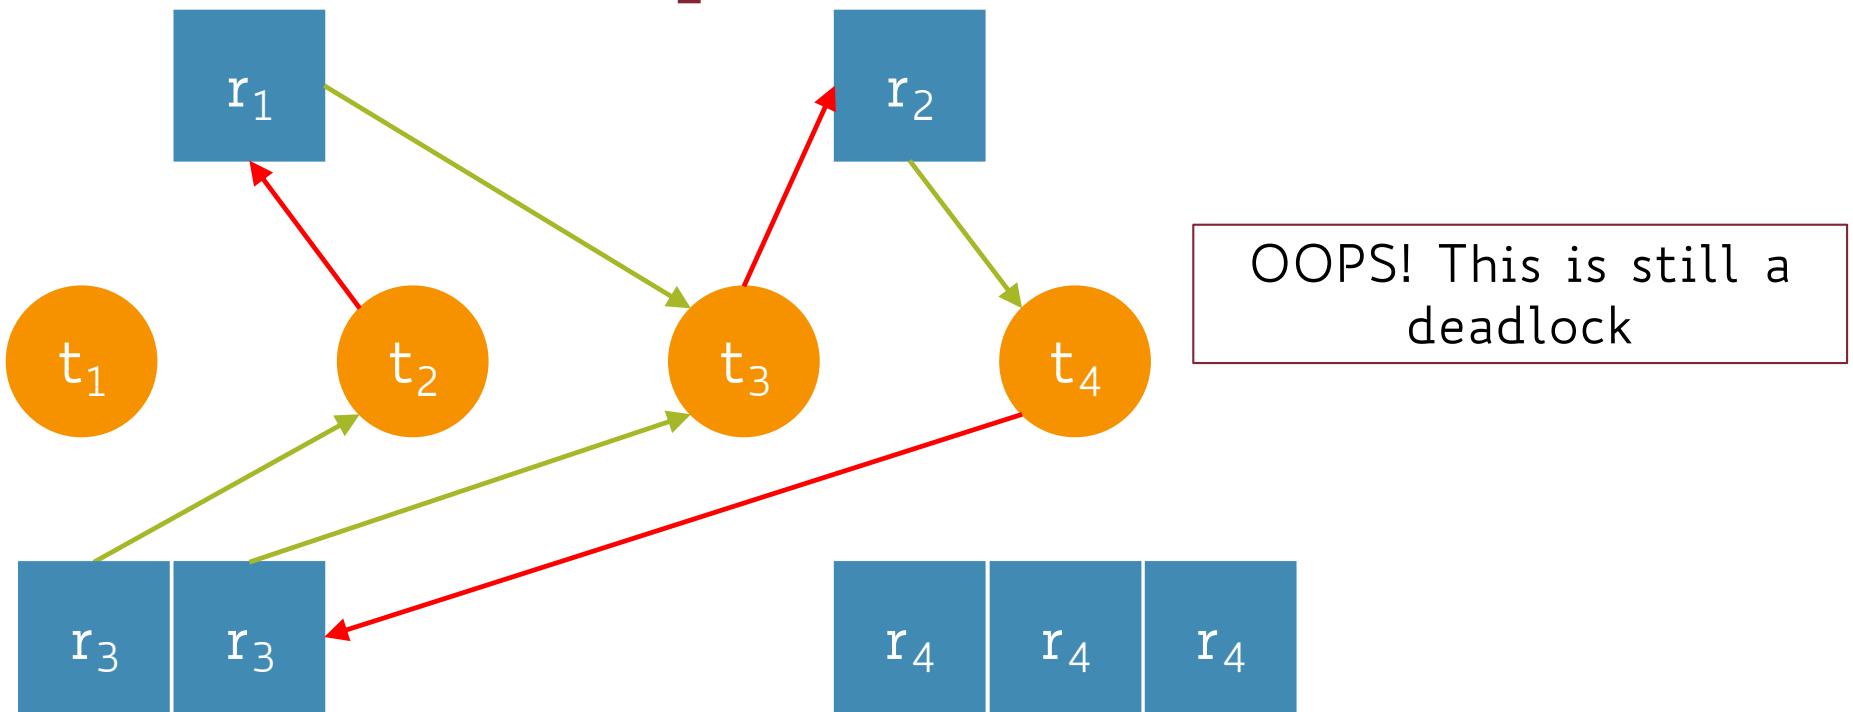
If the graph has cycles, deadlock **might** exist

We are assuming the **multiplicity** of each resource is **1** (i.e., one  $r_1$ , one  $r_2$ , etc.)


# Deadlock Detection: Resource Allocation Graph



# Deadlock Detection: Resource Allocation Graph



# Deadlock Detection: Resource Allocation Graph



# Deadlock: Detect and Correct It!

- Scan the Resource Allocation Graph (RAG) for cycles, and then break those!

# Deadlock: Detect and Correct It!

- Scan the Resource Allocation Graph (RAG) for cycles, and then break those!
- How? Several ways of doing it:

# Deadlock: Detect and Correct It!

- Scan the Resource Allocation Graph (RAG) for cycles, and then break those!
- How? Several ways of doing it:
  - Kill all the threads in the cycle (quite harsh, ugh?)

# Deadlock: Detect and Correct It!

- Scan the Resource Allocation Graph (RAG) for cycles, and then break those!
- How? Several ways of doing it:
  - Kill all the threads in the cycle (quite harsh, ugh?)
  - Kill all the threads one at a time, forcing each one of them to release resource(s)

# Deadlock: Detect and Correct It!

- Scan the Resource Allocation Graph (RAG) for cycles, and then break those!
- How? Several ways of doing it:
  - Kill all the threads in the cycle (quite harsh, ugh?)
  - Kill all the threads one at a time, forcing each one of them to release resource(s)
  - Preempt resources one at a time rolling back to a consistent status (e.g., common in database transactions)

# Deadlock: Detect and Correct It!

- Scan the Resource Allocation Graph (RAG) for cycles, and then break those!
- How? Several ways of doing it:
  - Kill all the threads in the cycle (quite harsh, ugh?)
  - Kill all the threads one at a time, forcing each one of them to release resource(s)
  - Preempt resources one at a time rolling back to a consistent status (e.g., common in database transactions)
- We would like to be more precise than that...

# Deadlock: Detect and Correct It!

- Detecting cycles on a directed graph  $G=(V, E)$  is a quite costly operation

# Deadlock: Detect and Correct It!

- Detecting cycles on a directed graph  $G=(V, E)$  is a quite costly operation
- Known algorithms based on **depth-first search (DFS)** take  $O(|V|+|E|)$  time

# Deadlock: Detect and Correct It!

- Detecting cycles on a directed graph  $G=(V, E)$  is a quite costly operation
- Known algorithms based on **depth-first search (DFS)** take  $O(|V|+|E|)$  time
- $O(|V|+|E|) \sim O(|V|^2)$  as  $|E| = O(|V|^2)$  for dense graphs, and  $|V| = \#\text{threads} + \#\text{resources}$

# Deadlock: Detect and Correct It!

- When to run such a detection algorithm?

# Deadlock: Detect and Correct It!

- When to run such a detection algorithm?
  - Before granting a resource → each granted request will take  $O(|V|^2)$

# Deadlock: Detect and Correct It!

- When to run such a detection algorithm?
  - Before granting a resource → each granted request will take  $O(|V|^2)$
  - When a request cannot be fulfilled → each failed request will take  $O(|V|^2)$

# Deadlock: Detect and Correct It!

- When to run such a detection algorithm?
  - Before granting a resource → each granted request will take  $O(|V|^2)$
  - When a request cannot be fulfilled → each failed request will take  $O(|V|^2)$
  - On a regular schedule or when the CPU is under-utilized

# Summary

- Deadlock → a situation in which a set of threads/processes cannot proceed because each one requires resources held by another

# Summary

- **Deadlock** → a situation in which a set of threads/processes cannot proceed because each one requires resources held by another
- **Prevention** → design resource allocation protocols that guarantee at least one of the 4 necessary deadlock conditions never holds

# Summary

- **Deadlock** → a situation in which a set of threads/processes cannot proceed because each one requires resources held by another
- **Prevention** → design resource allocation protocols that guarantee at least one of the 4 necessary deadlock conditions never holds
- **Avoidance** → scheduling threads so as to avoid deadlock

# Summary

- **Deadlock** → a situation in which a set of threads/processes cannot proceed because each one requires resources held by another
- **Prevention** → design resource allocation protocols that guarantee at least one of the 4 necessary deadlock conditions never holds
- **Avoidance** → scheduling threads so as to avoid deadlock
- **Detection and Recovery** → recognize deadlock after it has occurred and break it

# Summary

In practice, most OSs don't do anything and leave it all to applications



# Summary

After all, if deadlocks are rare, a non-solution like a hard reboot is often the best!

