

Systems and Networking I

Applied Computer Science and Artificial Intelligence
2023-2024

Gabriele Tolomei

Dipartimento di Informatica
Sapienza Università di Roma

tolomei@di.uniroma1.it



SAPIENZA
UNIVERSITÀ DI ROMA

LRU: Implementation Details

How could we implement LRU page replacement algorithm?

LRU: Implementation Details

How could we implement LRU page replacement algorithm?



First Idea

Keep a timestamp for each page with the time it has been last accessed
Remove the page with the highest difference w.r.t. current timestamp

LRU: Implementation Details

How could we implement LRU page replacement algorithm?



First Idea

Keep a timestamp for each page with the time it has been last accessed
Remove the page with the highest difference w.r.t. current timestamp

Problems?

LRU: Implementation Details

How could we implement LRU page replacement algorithm?



First Idea

Keep a timestamp for each page with the time it has been last accessed
Remove the page with the highest difference w.r.t. current timestamp

Problems?

Every time a page is
accessed its timestamp
must be updated

LRU: Implementation Details

How could we implement LRU page replacement algorithm?



First Idea

Keep a timestamp for each page with the time it has been last accessed
Remove the page with the highest difference w.r.t. current timestamp

Problems?

Every time a page is accessed its timestamp must be updated

Linear scan of all the pages to select the one to be removed

LRU: Implementation Details

How could we implement LRU page replacement algorithm?



Second Idea

Keep a list of pages with the most recently used in front and the least recently used at the end: every time a page is accessed move it to front

LRU: Implementation Details

How could we implement LRU page replacement algorithm?



Second Idea

Keep a list of pages with the most recently used in front and the least recently used at the end: every time a page is accessed move it to front

Problems?

LRU: Implementation Details

How could we implement LRU page replacement algorithm?



Second Idea

Keep a list of pages with the most recently used in front and the least recently used at the end: every time a page is accessed move it to front

Problems?



Still too expensive as the OS must change multiple pointers on each memory access

LRU: *Approximated* Implementation

- In practice, no need for perfect LRU

LRU: Approximated Implementation

- In practice, no need for perfect LRU
- Many systems provide some HW support which is enough to approximate LRU fairly well

LRU: Approximated Implementation

- In practice, no need for perfect LRU
- Many systems provide some HW support which is enough to approximate LRU fairly well
- **Single-Reference Bit** → 1 bit for each page table entry

LRU: Approximated Implementation

- In practice, no need for perfect LRU
- Many systems provide some HW support which is enough to approximate LRU fairly well
- **Single-Reference Bit** → 1 bit for each page table entry
 - Initially, all bits for all pages are set to 0

LRU: Approximated Implementation

- In practice, no need for perfect LRU
- Many systems provide some HW support which is enough to approximate LRU fairly well
- **Single-Reference Bit** → 1 bit for each page table entry
 - Initially, all bits for all pages are set to 0
 - On each access to a page, the HW sets the reference bit to 1

LRU: Approximated Implementation

- In practice, no need for perfect LRU
- Many systems provide some HW support which is enough to approximate LRU fairly well
- **Single-Reference Bit** → 1 bit for each page table entry
 - Initially, all bits for all pages are set to 0
 - On each access to a page, the HW sets the reference bit to 1
 - Enough to distinguish pages that have been accessed since the last clear

LRU: Approximated Implementation

- In practice, no need for perfect LRU
- Many systems provide some HW support which is enough to approximate LRU fairly well
- **Single-Reference Bit** → 1 bit for each page table entry
 - Initially, all bits for all pages are set to 0
 - On each access to a page, the HW sets the reference bit to 1
 - Enough to distinguish pages that have been accessed since the last clear
 - No total order of page access

LRU: Approximated Implementation

- Additional-Reference-Bits → E.g., 8 bits for each page table entry

LRU: Approximated Implementation

- *Additional-Reference-Bits* → E.g., 8 bits for each page table entry
 - At periodic intervals (clock interrupts), the OS takes over and right-shifts each of the reference bytes by one bit

LRU: Approximated Implementation

- **Additional-Reference-Bits** → E.g., 8 bits for each page table entry
 - At periodic intervals (clock interrupts), the OS takes over and right-shifts each of the reference bytes by one bit
 - The high-order (leftmost) bit is then filled in with the current value of the reference bit

LRU: Approximated Implementation

- **Additional-Reference-Bits** → E.g., 8 bits for each page table entry
 - At periodic intervals (clock interrupts), the OS takes over and right-shifts each of the reference bytes by one bit
 - The high-order (leftmost) bit is then filled in with the current value of the reference bit
 - At any given time, the page with the smallest value for the reference byte is the LRU page

LRU: Approximated Implementation

- **Additional-Reference-Bits** → E.g., 8 bits for each page table entry
 - At periodic intervals (clock interrupts), the OS takes over and right-shifts each of the reference bytes by one bit
 - The high-order (leftmost) bit is then filled in with the current value of the reference bit
 - At any given time, the page with the smallest value for the reference byte is the LRU page
- The specific number of bits used and the frequency with which the reference byte is updated are adjustable

LRU: Approximated Implementation

- Second Chance Algorithm → Single-Reference Bit + FIFO

LRU: Approximated Implementation

- Second Chance Algorithm → Single-Reference Bit + FIFO
- OS keeps frames in a FIFO circular list

LRU: Approximated Implementation



- Second Chance Algorithm → Single-Reference Bit + FIFO
- OS keeps frames in a FIFO circular list
- On every memory access, the reference bit is set to 1

LRU: Approximated Implementation


- Second Chance Algorithm → Single-Reference Bit + FIFO
- OS keeps frames in a FIFO circular list
- On every memory access, the reference bit is set to 1
- On a page fault, the OS scans the list of frames, checking the reference bit of the frame:
 - If this is 0, it replaces the page and sets it to 1
 - If this is 1, it sets it to 0 (second chance) and move to the next frame

Second Chance Algorithm (Clock)

A raw partitioning into: young vs. old frames



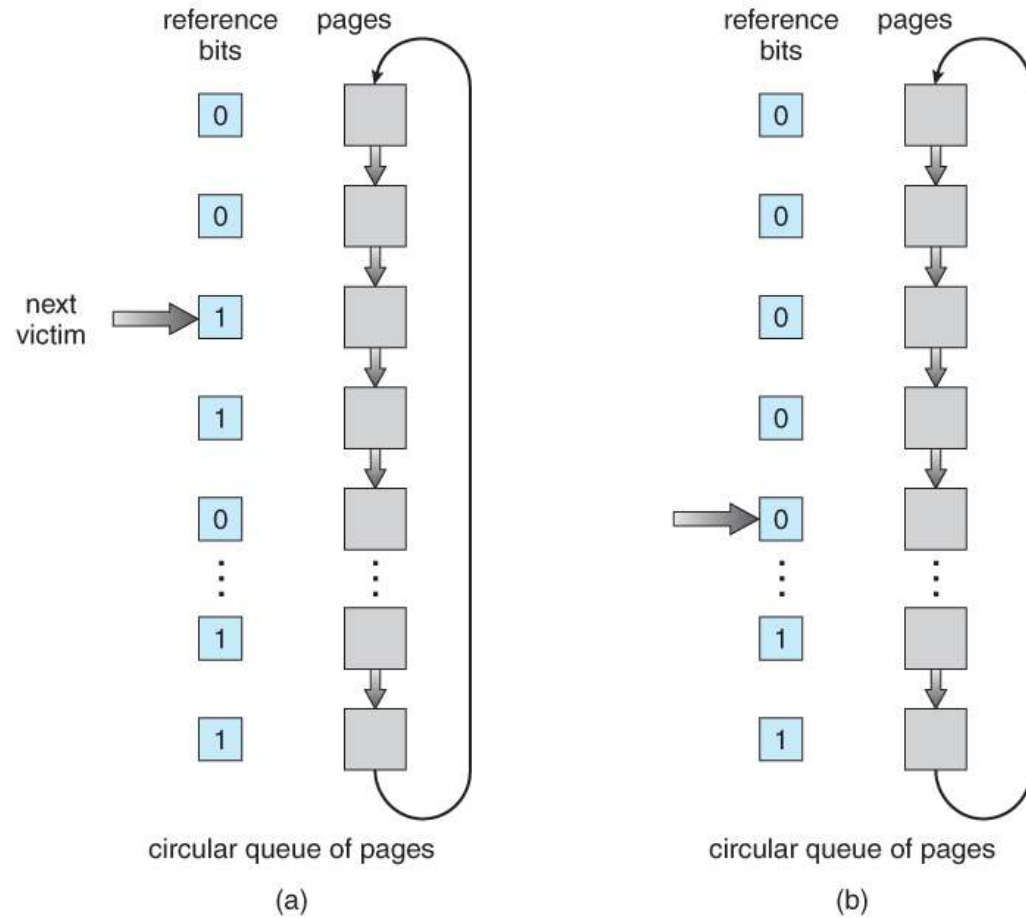
Second Chance Algorithm (Clock)



A raw partitioning into: young vs. old frames

Less accurate than additional-reference-bits

Second Chance Algorithm (Clock)




A raw partitioning into: young vs. old frames

Less accurate than additional-reference-bits

Fast, since it needs only to set a single bit on each memory reference (no need for a shift)

Second Chance Algorithm (Clock)



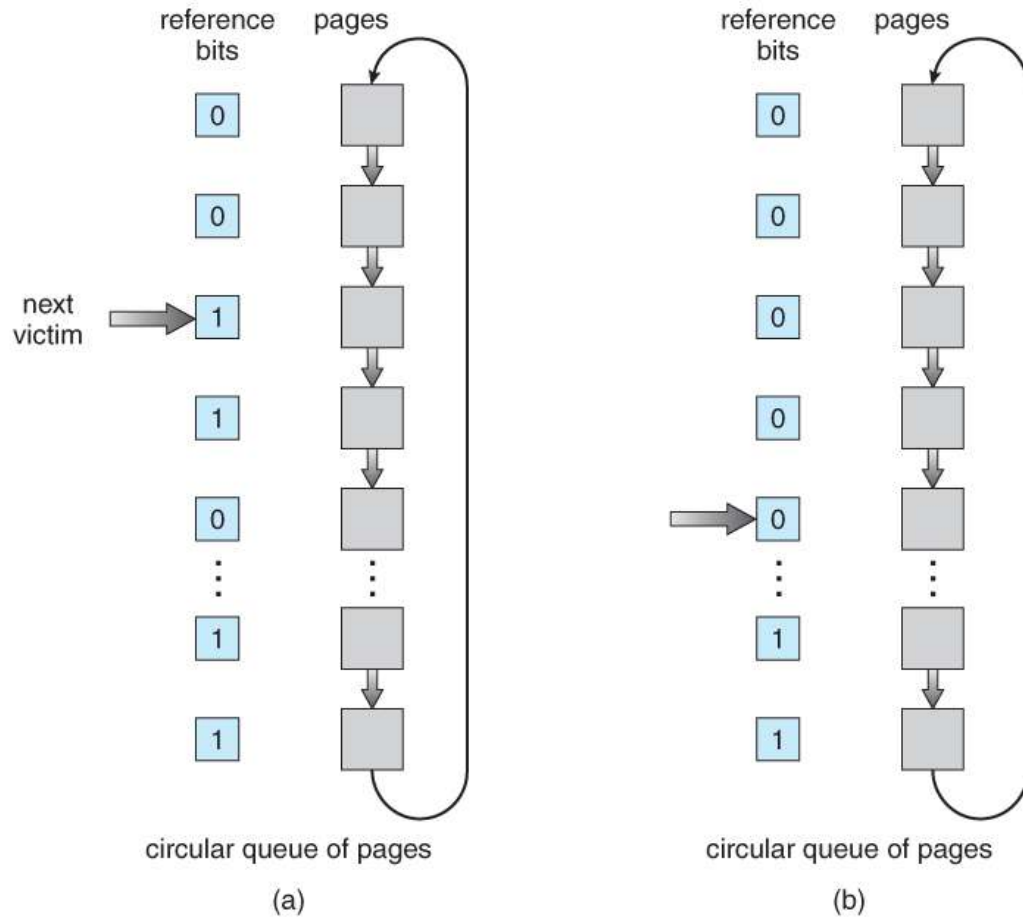
A raw partitioning into: young vs. old frames

Less accurate than additional-reference-bits

Fast, since it needs only to set a single bit on each memory reference (no need for a shift)

Page fault management is quicker as there is no need to scan the whole list of frames (on average) unless every frame has its bit set

Second Chance Algorithm (Clock)



A raw partitioning into: young vs. old frames

Less accurate than additional-reference-bits

Fast, since it needs only to set a single bit on each memory reference (no need for a shift)

Page fault management is quicker as there is no need to scan the whole list of frames (on average) unless every frame has its bit set

This algorithm is also known as **clock** because it mimics the hands of a clock

Enhanced Second Chance Algorithm

- Page replacement generally involves 2 I/O operations:
 - write the evicted page back to disk
 - read the newly referenced page from disk

Enhanced Second Chance Algorithm

- Page replacement generally involves 2 I/O operations:
 - write the evicted page back to disk
 - read the newly referenced page from disk
- **Intuition:** It is cheaper to replace a page which has not been modified, since the OS does not need to write this back to disk

Enhanced Second Chance Algorithm

- OS should give preference to paging-out un-modified frames
- Yet, it can proactively write to disk modified frames for later

Enhanced Second Chance Algorithm

- HW keeps a modify bit (in addition to the reference bit)
 - 1 means the page has been modified (different from the copy on disk)
 - 0 means the page is the same as the one stored on disk

Enhanced Second Chance Algorithm

- HW keeps a modify bit (in addition to the reference bit)
 - 1 means the page has been modified (different from the copy on disk)
 - 0 means the page is the same as the one stored on disk
- Use both the reference and modify bits (r, m) to classify pages into:
 - (0, 0): neither recently used nor modified;
 - (0, 1): not recently used, but modified;
 - (1, 0): recently used, but clean
 - (1, 1): recently used and modified

Enhanced Second Chance Algorithm

- This algorithm searches the page table in a circular fashion as before, yet making 4 distinct passes (one for each category)

Enhanced Second Chance Algorithm

- This algorithm searches the page table in a circular fashion as before, yet making 4 distinct passes (one for each category)
- First, it looks for a (0, 0) frame, and if it can't find one, it makes another pass looking for a (0, 1), etc.

Enhanced Second Chance Algorithm

- This algorithm searches the page table in a circular fashion as before, yet making 4 distinct passes (one for each category)
- First, it looks for a (0, 0) frame, and if it can't find one, it makes another pass looking for a (0, 1), etc.
- The first page with the lowest numbered category is replaced

Enhanced Second Chance Algorithm

- This algorithm searches the page table in a circular fashion as before, yet making 4 distinct passes (one for each category)
- First, it looks for a (0, 0) frame, and if it can't find one, it makes another pass looking for a (0, 1), etc.
- The first page with the lowest numbered category is replaced
- Prioritize replacement of clean pages if possible

Multiprogramming and Thrashing

- So far, we have implicitly assumed a single process is on the system

Multiprogramming and Thrashing

- So far, we have implicitly assumed a single process is on the system
- Multiple processes can however run concurrently on a single-CPU system

Multiprogramming and Thrashing

- So far, we have implicitly assumed a single process is on the system
- Multiple processes can however run concurrently on a single-CPU system
- The degree of multiprogramming is not fixed apriori, yet it is driven by the locality reference (a.k.a. 90÷10 rule)

Multiprogramming and Thrashing

- So far, we have implicitly assumed a single process is on the system
- Multiple processes can however run concurrently on a single-CPU system
- The degree of multiprogramming is not fixed apriori, yet it is driven by the locality reference (a.k.a. 90÷10 rule)
- This allows a system to load the **working set** (i.e., few pages) of many processes, thereby increasing the degree of multiprogramming


Multiprogramming and Thrashing

- When the degree of multiprogramming is too high, active working sets of running processes may saturate the whole memory capacity

Multiprogramming and Thrashing


- When the degree of multiprogramming is too high, active working sets of running processes may saturate the whole memory capacity
- **Thrashing** → Memory is over-committed and pages are continuously tossed out while they are still in use
 - Memory access time approaches disk access time due to many page faults
 - Drastic degradation of performance

Multiprogramming and Thrashing



CPU utilization drops after a certain degree of multiprogramming


Multiprogramming and Thrashing



CPU utilization drops after a certain degree of multiprogramming

Eventually, also CPU-bound processes turn into I/O-bound ones (due to page faults)

Multiprogramming and Thrashing




CPU utilization drops after a certain degree of multiprogramming

Eventually, also CPU-bound processes turn into I/O-bound ones (due to page faults)

What can we do to limit thrashing in a multiprogrammed system?

Multiprogramming and Thrashing



CPU utilization drops after a certain degree of multiprogramming

Eventually, also CPU-bound processes turn into I/O-bound ones (due to page faults)

What can we do to limit thrashing in a multiprogrammed system?

Fixing the degree of multi-programming apriori may be a too inflexible option

Allocation/Replacement Policies

Ultimately, we want to give each process enough memory so as to avoid thrashing

Allocation/Replacement Policies

Ultimately, we want to give each process enough memory so as to avoid thrashing

Global Allocation/Replacement

- All pages from all processes are in a single pool (single LRU queue)
- Upon page replacement, any page may be a potential victim, whether it currently belongs to the process seeking a free frame or not
- **PRO:** flexibility
- **CON:** thrashing more likely (no isolation)

Allocation/Replacement Policies

Ultimately, we want to give each process enough memory so as to avoid thrashing

Local Allocation/Replacement

- Each process has its own fixed pool of frames
- Run only group of processes that fits in memory
- LRU replacement affects only each process' frames
- **PRO:** isolation
- **CON:** performance (a process may not be given enough memory)

Allocation/Replacement Policies

Ultimately, we want to give each process enough memory so as to avoid thrashing

Global Allocation/Replacement

- All pages from all processes are in a single pool (single LRU queue)
- Upon page replacement, any page may be a potential victim, whether it currently belongs to the process seeking a free frame or not
- **PRO:** flexibility
- **CON:** thrashing more likely (no isolation)

12/20/23

Local Allocation/Replacement

- Each process has its own fixed pool of frames
- Run only group of processes that fits in memory
- LRU replacement affects only each process' frames
- **PRO:** isolation
- **CON:** performance (a process may not be given enough memory)

53

Local Allocation/Replacement

m = number of available physical page frames

n = number of processes

S_i = size of the i -th process; $S = \sum_{i=1}^n S_i$ = total size of all processes

Equal Allocation/Replacement: $\frac{m}{n}$

Proportional Allocation/Replacement: $\frac{m * S_i}{S}$

Local Allocation/Replacement

m = number of available physical page frames

n = number of processes

S_i = size of the i -th process; $S = \sum_{i=1}^n S_i$ = total size of all processes

Equal Allocation/Replacement: $\frac{m}{n}$

Proportional Allocation/Replacement: $\frac{m * S_i}{S}$

Variations on proportional allocation could consider priority of process rather than just their size

Local Allocation/Replacement

m = number of available physical page frames

n = number of processes

S_i = size of the i -th process; $S = \sum_{i=1}^n S_i$ = total size of all processes

Equal Allocation/Replacement: $\frac{m}{n}$

Proportional Allocation/Replacement: $\frac{m * S_i}{S}$

Variations on proportional allocation could consider priority of process rather than just their size

As allocations fluctuate over time, so does m
(processes must be swapped out or not started if not enough frames)

Proportional Allocation/Replacement

- This implicitly assumes that a large process will also refer to a large amount of memory

Proportional Allocation/Replacement

- This implicitly assumes that a large process will also refer to a large amount of memory
- Intuitively this makes sense: the more memory a process needs the higher the chance it will refer a much larger memory range

Proportional Allocation/Replacement

- This implicitly assumes that a large process will also refer to a large amount of memory
- Intuitively this makes sense: the more memory a process needs the higher the chance it will refer a much larger memory range
- However, there might be cases where this is not true
 - e.g., a process allocates a 1GB array but only uses a small portion of it

Proportional Allocation/Replacement

- This implicitly assumes that a large process will also refer to a large amount of memory
- Intuitively this makes sense: the more memory a process needs the higher the chance it will refer a much larger memory range
- However, there might be cases where this is not true
 - e.g., a process allocates a 1GB array but only uses a small portion of it
- In other words, the working set of a process may not be correlated with its memory footprint


Matching the Working Set

- **Goal** → Give each process enough frames to contain its working set
 - Informally, the working set is the set of pages the process is using "right now"
 - More formally, it is the set of all pages that the process has referenced during the past T units of time (e.g., seconds)

Matching the Working Set


- How does the OS pick T?
 - 1 page fault takes order of 10 msec to be served
 - 10 msec ~ 10 million instructions
 - T needs to account for a lot more than 10 million instructions

Determining the Working Set



The selection of Δ is critical to the success of the working set model


Determining the Working Set



The selection of Δ is critical to the success of the working set model

Δ too small
it does not encompass all of the
pages of the current locality


Determining the Working Set



The selection of Δ is critical to the success of the working set model

Δ too large
it includes pages that are no longer frequently accessed

Determining the Working Set



The selection of Δ is critical to the success of the working set model

Δ too small

it does not encompass all of the pages of the current locality

Δ too large

it includes pages that are no longer frequently accessed

Exact tracking is expensive: update the working set at each memory access

Approximating the Working Set

- Computing the working set exactly requires to keep track of a moving window of size Δ

Approximating the Working Set

- Computing the working set exactly requires to keep track of a moving window of size Δ
- At each memory reference a new reference appears at one end and the oldest reference drops off the other end

Approximating the Working Set

- Computing the working set exactly requires to keep track of a moving window of size Δ
- At each memory reference a new reference appears at one end and the oldest reference drops off the other end
- To avoid the overhead of keeping a list of the last Δ referenced pages, the working set is often implemented with **sampling**

Approximating the Working Set


- Computing the working set exactly requires to keep track of a moving window of size Δ
- At each memory reference a new reference appears at one end and the oldest reference drops off the other end
- To avoid the overhead of keeping a list of the last Δ referenced pages, the working set is often implemented with **sampling**
- Every k memory references (e.g., $k = 1,000$), consider the working set to be all pages referenced within *that* period of time

Tracking Page Fault Rate

- Ultimately, our goal is to minimize the **page fault rate**

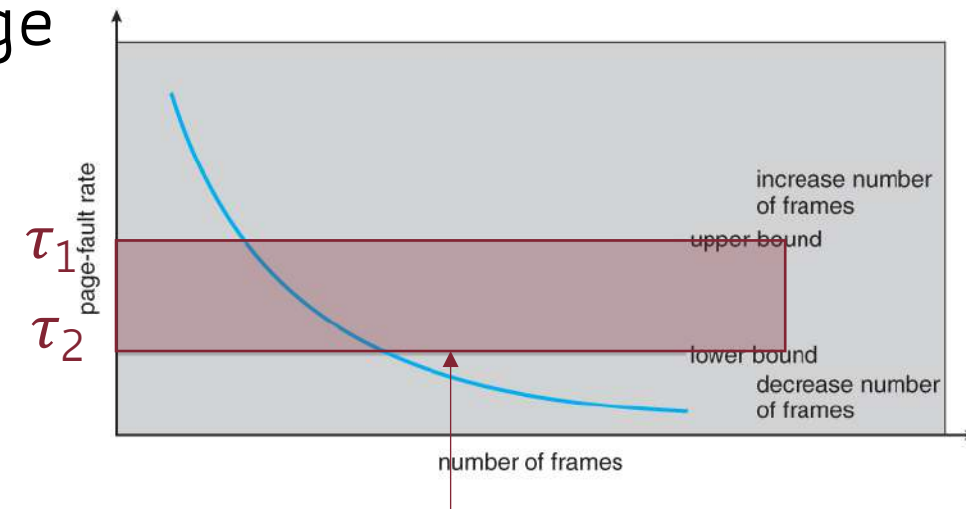
Tracking Page Fault Rate

- Ultimately, our goal is to minimize the **page fault rate**
- A more direct approach is to track page fault rate of each process:
 - If the page fault rate is above a given threshold $\tau_1 \rightarrow$ give it more frames
 - If the page fault rate is below a given threshold $\tau_2 \rightarrow$ reclaim some frames



Tracking Page Fault Rate

- Ultimately, our goal is to minimize the **page fault rate**
- A more direct approach is to track page fault rate of each process:
 - If the page fault rate is above a given threshold $\tau_1 \rightarrow$ give it more frames
 - If the page fault rate is below a given threshold $\tau_2 \rightarrow$ reclaim some frames



Dynamically adjust allocated frames so as to keep processes in this area

Kernel Memory

- So far, we only considered memory allocation for user processes

Kernel Memory

- So far, we only considered memory allocation for user processes
- But kernel needs memory to store things too: code and data structures like PCB, page tables, etc.

Kernel Memory

- So far, we only considered memory allocation for user processes
- But kernel needs memory to store things too: code and data structures like PCB, page tables, etc.
- Kernel does not use any of the advanced mechanisms seen so far
 - No paging → what if a page fault occurs for the kernel?

Kernel Memory: Buddy Allocator

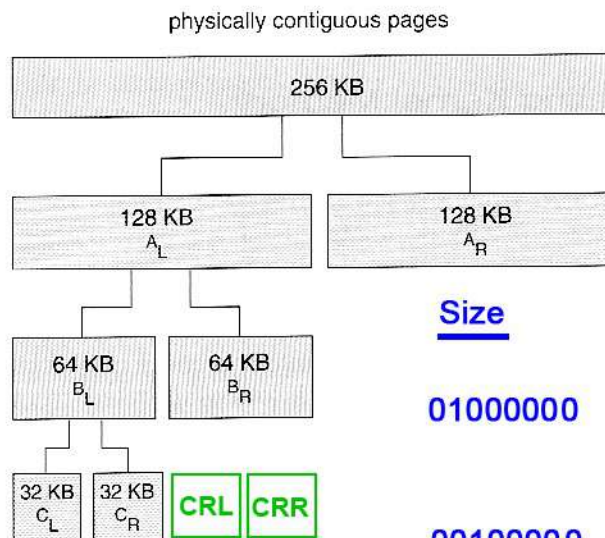


Figure 9.27 Buddy system allocation.

Buddy Addresses

00000000

00000000 10000000

Size

01000000

00000000 01000000

00100000


00000000 00100000

01000000

01100000

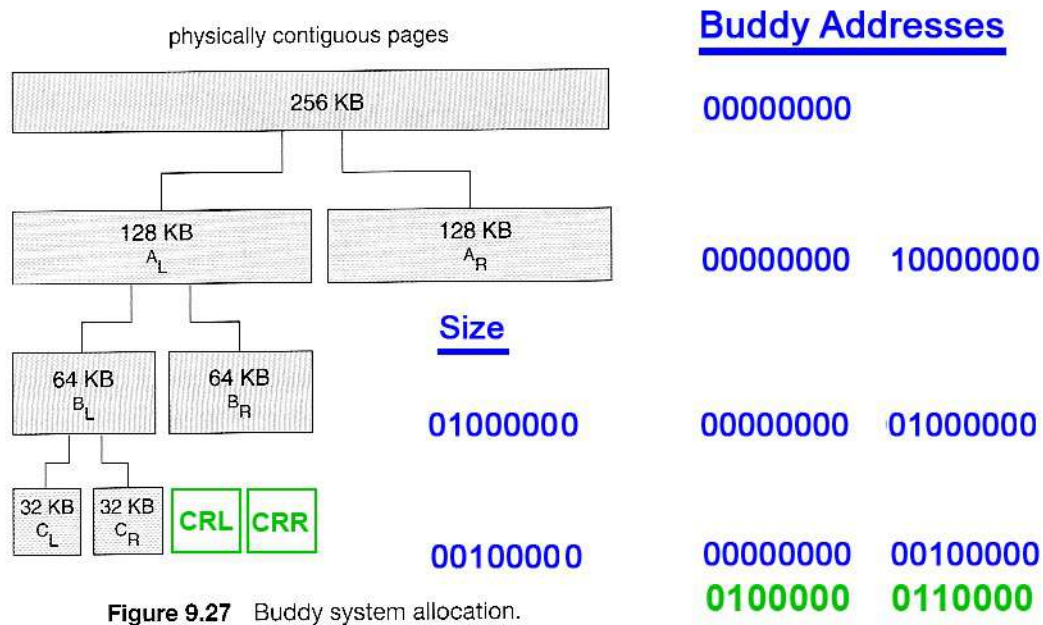
- Allocates memory using a power of 2 allocator (e.g., 4KiB, 8KiB, 16KiB), rounding up to the next nearest power of two if necessary

Kernel Memory: Buddy Allocator




- Allocates memory using a power of 2 allocator (e.g., 4KiB, 8KiB, 16KiB), rounding up to the next nearest power of two if necessary
- If a block of the correct size is not available, then one is formed by (repeatedly) splitting the next larger block in two

Kernel Memory: Buddy Allocator



- Allocates memory using a power of 2 allocator (e.g., 4KiB, 8KiB, 16KiB), rounding up to the next nearest power of two if necessary
- If a block of the correct size is not available, then one is formed by (repeatedly) splitting the next larger block in two
- Can lead to internal fragmentation

Kernel Memory: Slab Allocator



- Group of objects of the same size in a **slab**
- Object cache points to one or more slabs
- Separate cache for each kernel data structure (e.g., PCB)
- No internal fragmentation
- Used in Solaris and Linux

Page Sizes

- Reasons for **small** pages?

Page Sizes

- Reasons for **small** pages?
 - Decreasing internal fragmentation
 - Higher degree of multiprogramming

Page Sizes

- Reasons for **small** pages?
 - Decreasing internal fragmentation
 - Higher degree of multiprogramming
- Reasons for **large** pages?

Page Sizes

- Reasons for **small** pages?
 - Decreasing internal fragmentation
 - Higher degree of multiprogramming
- Reasons for **large** pages?
 - Smaller page table size (i.e., smaller number of page table entries)
 - Fewer page faults (locality reference)
 - Amortizes disk overhead (reading a 1KiB page from disk takes approximately the same as reading an 8KiB one)

Summary of Page Replacement

- The choice of page replacement algorithm is crucial when physical memory is limited
 - All algorithms approach to the optimum as the physical memory allocated to a process approaches to the virtual memory size

Summary of Page Replacement

- The choice of page replacement algorithm is crucial when physical memory is limited
 - All algorithms approach to the optimum as the physical memory allocated to a process approaches to the virtual memory size
- The more processes running concurrently, the less physical memory each one can have

Summary of Page Replacement

- The choice of page replacement algorithm is crucial when physical memory is limited
 - All algorithms approach to the optimum as the physical memory allocated to a process approaches to the virtual memory size
- The more processes running concurrently, the less physical memory each one can have
- The OS must choose how many processes (and the number of frames per process) can share memory